

## Linear Regression Model

```
1 X = led_df.drop('BMI', axis=1)
2 y = led_df['BMI']
```

```
1 print("X", X.shape, "y", y.shape)
```

```
X (2938, 21)
y= (2938,)
```

## Train Test Splitting

```
1 !pip install scikit-learn
```

```
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-packages (1.2.2)
Requirement already satisfied: numpy>=1.17.3 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.25.2)
Requirement already satisfied: scipy>=1.3.2 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.11.4)
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.4.0)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (3.4.0)
```

```
1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 import hvplot.pandas
6 %matplotlib inline
7
8 from sklearn.model_selection import train_test_split
9 from sklearn import metrics
10 from sklearn.linear_model import LinearRegression
```

```
1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=101)
```

```
1 X_train.shape
(2056, 21)
```

```
1 X_test.shape
(882, 21)
```

```
1 model = LinearRegression()
1 model.fit(X_train, y_train)
```

```
▼ LinearRegression
  LinearRegression()
```

## ▼ Model Evaluation

```
1 model.coef_
array([
  7.57056771e-03, -4.63411687e-02, -2.07939848e+00,  5.56411501e-01,
 -6.24185007e-03, -3.41735872e-03, -1.14684826e-01,  3.06103000e-04,
 2.83703325e-02, -1.07978944e-04,  9.05328887e-03, -7.11044654e-04,
 3.56594950e-01, -1.75437562e-02,  1.36511969e-01, -8.45656970e-05,
 6.05489000e-09, -4.91879831e-01, -1.07257902e+00,  4.62048803e+00,
 1.01030665e+00])
```

```
1 pd.DataFrame(model.coef_, X.columns, columns=['Coefficients'])
```

	Coedicients	
Country	7.570568e-03	
Year	-4.634117e-02	
Status	-2.079398e+00	
Life expectancy	5.564115e-01	
Adult Mortality	-6.241850e-03	
infant deaths	-3.417359e-03	
Alcohol	-1.146848e-01	
percentage expenditure	3.061030e-04	
Hepatitis B	2.837033e-02	
Measles	-1.079789e-04	
under-five deaths	9.053289e-03	
Polio	-7.110447e-04	
Total expenditure	3.565949e-01	
Diphtheria	-1.754376e-02	
HIV/AIDS	1.365120e-01	
GDP	-8.456570e-05	
Population	6.054890e-09	
thinness 1-19 years	-4.918798e-01	
thinness 5-9 years	-1.072579e+00	
Income composition of resources	4.620488e+00	
Schooling	1.010307e+00	

## ▼ Predictions from our Model

```
1 y_pred = model.predict(X_test)
```

## ▼ Regression Evaluation Metrics

```
1 MAE = metrics.mean_absolute_error(y_test, y_pred)
2 MSE = metrics.mean_squared_error(y_test, y_pred)
3 RMSE = np.sqrt(MSE)
```

```
1 MAE
```

```
10.684448166792613
```

```
1 MSE
```

```
212.90158397376152
```

```
1 RMSE
```

```
14.591147452265758
```

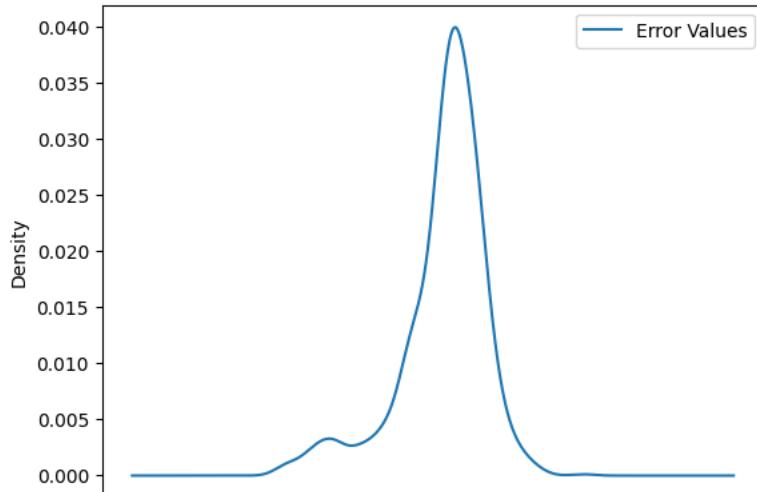
```
1 led_df['BMI'].mean()
```

```
38.02014976174268
```

## ▼ Residual Histogram

```
1 test_residual = y_test - y_pred
1 pd.DataFrame({'Error Values': (test_residual)}).plot.kde()
```

```
<Axes: ylabel='Density'>
```



```
1 sns.displot(test_residual, bins=25, kde=True)
```

```
<seaborn.axisgrid.FacetGrid at 0x7c501fc41990>
```

