

Seatwork 11.1 Exploratory Data Analysis for Machine Learning

Name: Cuadra, Audrick Zander G.

Section: CPE22S3

Date: April 22, 2024

Instructions:

- Download the datasets here:
- For Linear Regression Analysis: <https://archive-beta.ics.uci.edu/dataset/10/automobileLinks> to an external site.
- For Logistic Regression Analysis: <https://archive-beta.ics.uci.edu/dataset/109/wine>
- Perform exploratory data analysis (which must include data pre-processing/wrangling).
- Submit the notebook with the cleaned data and the EDA.

Note:

- Your submission must be PDF file.
- However, submit the link of your python notebook and submit the link in the comments.

Automobile

```
1 # Install the ucimlrepo package
2 !pip install ucimlrepo

Requirement already satisfied: ucimlrepo in /usr/local/lib/python3.10/dist-packages (0.0.6)

1 # Import the dataset into your code
2 from ucimlrepo import fetch_ucirepo
3
4 # fetch dataset
5 automobile = fetch_ucirepo(id=10)
6
7 # data (as pandas dataframes)
8 X = automobile.data.features
9 y = automobile.data.targets
10
11 # metadata
12 print(automobile.metadata)
13
14 # variable information
15 print(automobile.variables)
```

	name	role	type	demographic	\
0	price	Feature	Continuous	None	
1	highway-mpg	Feature	Continuous	None	
2	city-mpg	Feature	Continuous	None	
3	peak-rpm	Feature	Continuous	None	
4	horsepower	Feature	Continuous	None	
5	compression-ratio	Feature	Continuous	None	
6	stroke	Feature	Continuous	None	
7	bore	Feature	Continuous	None	
8	fuel-system	Feature	Categorical	None	
9	engine-size	Feature	Continuous	None	
10	num-of-cylinders	Feature	Integer	None	
11	engine-type	Feature	Categorical	None	
12	curb-weight	Feature	Continuous	None	
13	height	Feature	Continuous	None	
14	width	Feature	Continuous	None	
15	length	Feature	Continuous	None	
16	wheel-base	Feature	Continuous	None	
17	engine-location	Feature	Binary	None	
18	drive-wheels	Feature	Categorical	None	
19	body-style	Feature	Categorical	None	
20	num-of-doors	Feature	Integer	None	
21	aspiration	Feature	Binary	None	
22	fuel-type	Feature	Binary	None	
23	make	Feature	Categorical	None	
24	normalized-losses	Feature	Continuous	None	
25	symboling	Target	Integer	None	

	description	units	missing_values
0	continuous from 5118 to 45400	None	yes



```
1          continuous from 16 to 54 None      no
2          continuous from 13 to 49 None      no
3          continuous from 4150 to 6600 None    yes
4          continuous from 48 to 288 None      yes
5          continuous from 7 to 23 None      no
6          continuous from 2.07 to 4.17 None    yes
7          continuous from 2.54 to 3.94 None    yes
8      1bbl, 2bbl, 4bbl, idi, mfi, mpfi, spdi, spfi None no
9          continuous from 61 to 326 None      no
10         eight, five, four, six, three, twelve, two None no
11         dohc, dohcvt, l, ohc, ohcf, ohcv, rotor None no
12         continuous from 1488 to 4066 None    no
13         continuous from 47.8 to 59.8 None    no
14         continuous from 60.3 to 72.3 None    no
15         continuous from 141.1 to 208.1 None   no
16         continuous from 86.6 120.9 None      no
17         front, rear None                    no
18         4wd, fwd, rwd None                  no
19         hardtop, wagon, sedan, hatchback, convertible None no
20         four, two None                      yes
21         std, turbo None                     no
22         diesel, gas None                    no
23     alfa-romero, audi, bmw, chevrolet, dodge, hond... None no
24         continuous from 65 to 256 None      yes
25         -3, -2, -1, 0, 1, 2, 3 None        no
```

```
1 # showing a sample of the 'X' dataframe
2 X.head()
```

	price	highway- mpg	city- mpg	peak- rpm	horsepower	compression- ratio	stroke	bore	fuel- system	er
0	13495.0	27	21	5000.0	111.0	9.0	2.68	3.47	mpfi	
1	16500.0	27	21	5000.0	111.0	9.0	2.68	3.47	mpfi	
2	16500.0	26	19	5000.0	154.0	9.0	3.47	2.68	mpfi	
3	13950.0	30	24	5500.0	102.0	10.0	3.40	3.19	mpfi	
4	17450.0	22	18	5500.0	115.0	8.0	3.40	3.19	mpfi	

5 rows × 25 columns

```
1 # showing a sample of the 'y' dataframe
2 y.head()
```

	symboling	
0	3	
1	3	
2	1	
3	2	
4	2	

Next steps:

☒

[View recommended plots](#)

```
1 # importing all the necessary libraries
2 import pandas as pd
3 import numpy as np
4
5 # concatenating the 'X' and 'y' dataframe
6 autom_df = pd.concat([X, y], axis=1)
7 autom_df
```

	price	highway- mpg	city- mpg	peak- rpm	horsepower	compression- ratio	stroke	bore	fuel- system
0	13495.0	27	21	5000.0	111.0	9.0	2.68	3.47	mpfi
1	16500.0	27	21	5000.0	111.0	9.0	2.68	3.47	mpfi
2	16500.0	26	19	5000.0	154.0	9.0	3.47	2.68	mpfi
3	13950.0	30	24	5500.0	102.0	10.0	3.40	3.19	mpfi
4	17450.0	22	18	5500.0	115.0	8.0	3.40	3.19	mpfi
...
200	16845.0	28	23	5400.0	114.0	9.5	3.15	3.78	mpfi
201	19045.0	25	19	5300.0	160.0	8.7	3.15	3.78	mpfi
202	21485.0	23	18	5500.0	134.0	8.8	2.87	3.58	mpfi
203	22470.0	27	26	4800.0	106.0	23.0	3.40	3.01	idi
204	22625.0	25	19	5400.0	114.0	9.5	3.15	3.78	mpfi

205 rows × 26 columns

```

1 # creating a function that detects and displays the number of duplicates in the dataframe
2 def countDuplicate(data):
3     if data.duplicated().any():
4         count = data.duplicated().sum()
5         print(count)
6     else:
7         return "No Duplicates Found!"

```

```

1 # checking if the concatenated dataframe has any duplicates
2 countDuplicate(autom_df)

```

'No Duplicates Found!'

```

1 # checking the number of nulls in the dataframe
2 autom_df.isnull().sum()

```

```

price                4
highway-mpg          0
city-mpg              0
peak-rpm             2
horsepower           2
compression-ratio    0
stroke               4
bore                 4
fuel-system          0
engine-size          0
num-of-cylinders     0
engine-type          0
curb-weight          0
height              0
width               0
length              0
wheel-base          0
engine-location      0
drive-wheels         0
body-style           0
num-of-doors         2
aspiration           0
fuel-type            0
make                0
normalized-losses    41
symboling            0
dtype: int64

```

```
1 autom_df['price'].value_counts()
```

```

price
8921.0      2
18150.0     2
8845.0      2
8495.0      2
7609.0      2
..
45400.0     1

```

```

16503.0    1
5389.0     1
6189.0     1
22625.0    1
Name: count, Length: 186, dtype: int64

```

```
1 autom_df.isnull().sum()
```

```

price                4
highway-mpg          0
city-mpg             0
peak-rpm             2
horsepower           2
compression-ratio    0
stroke              4
bore                4
fuel-system          0
engine-size          0
num-of-cylinders     0
engine-type          0
curb-weight          0
height              0
width               0
length              0
wheel-base          0
engine-location      0
drive-wheels         0
body-style           0
num-of-doors         2
aspiration           0
fuel-type            0
make                0
normalized-losses    41
symboling            0
dtype: int64

```

```
1 # using mean to supply each columns that contains a null
```

```
2 autom_df['price'].fillna(autom_df['price'].median(), inplace=True)
```

```
1 autom_df['peak-rpm'].fillna(autom_df['peak-rpm'].median(), inplace=True)
```

```
1 autom_df['horsepower'].fillna(autom_df['horsepower'].median(), inplace=True)
```

```
1 autom_df['stroke'].fillna(autom_df['stroke'].median(), inplace=True)
```

```
1 autom_df['bore'].fillna(autom_df['bore'].median(), inplace=True)
```

```
1 autom_df['num-of-doors'].fillna(autom_df['num-of-doors'].median(), inplace=True)
```

```
1 autom_df['normalized-losses'].fillna(autom_df['normalized-losses'].median(), inplace=True)
```

```
1 # checking if there are still any nulls in the dataframe
```

```
2 autom_df.isnull().sum()
```

```

price                0
highway-mpg          0
city-mpg             0
peak-rpm             0
horsepower           0
compression-ratio    0
stroke              0
bore                0
fuel-system          0
engine-size          0
num-of-cylinders     0
engine-type          0
curb-weight          0
height              0
width               0
length              0
wheel-base          0
engine-location      0
drive-wheels         0
body-style           0
num-of-doors         0
aspiration           0
fuel-type            0
make                0
normalized-losses    0
symboling            0
dtype: int64

```

```
1 # checking all the categorical values
2 autom_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
#   Column                Non-Null Count  Dtype
---  -
0   price                  205 non-null   float64
1   highway-mpg            205 non-null   int64
2   city-mpg               205 non-null   int64
3   peak-rpm               205 non-null   float64
4   horsepower              205 non-null   float64
5   compression-ratio      205 non-null   float64
6   stroke                 205 non-null   float64
7   bore                   205 non-null   float64
8   fuel-system            205 non-null   object
9   engine-size            205 non-null   int64
10  num-of-cylinders        205 non-null   int64
11  engine-type            205 non-null   object
12  curb-weight            205 non-null   int64
13  height                 205 non-null   float64
14  width                  205 non-null   float64
15  length                 205 non-null   float64
16  wheel-base             205 non-null   float64
17  engine-location        205 non-null   object
18  drive-wheels           205 non-null   object
19  body-style             205 non-null   object
20  num-of-doors           205 non-null   float64
21  aspiration             205 non-null   object
22  fuel-type              205 non-null   object
23  make                   205 non-null   object
24  normalized-losses      205 non-null   float64
25  symboling              205 non-null   int64
dtypes: float64(12), int64(6), object(8)
memory usage: 41.8+ KB
```

```
1 # This is to allow the access for the dataframe in which the int conversion hasn't occurred
2 autom_cat = autom_df.copy()
```

```
1 # creating a function that converts objects into numerical values
2 def preprocessing(data, catlist):
3     if data[catlist].dtypes == 'object':
4         cat_val = data[catlist].unique()
5         range_val = range(1, len(cat_val)+1)
6         map = dict(zip(cat_val, range_val))
7         print(f"{catlist}:", map)
8         data[catlist] = data[catlist].map(map)
9     return data
10 for i in autom_df.select_dtypes(include=['object']).columns:
11     preprocessing(autom_df, i)

fuel-system: {'mpfi': 1, '2bbl': 2, 'mfi': 3, '1bbl': 4, 'spfi': 5, '4bbl': 6, 'idi': 7, 'spdi': 8}
engine-type: {'dohc': 1, 'ohcv': 2, 'ohc': 3, 'l': 4, 'rotor': 5, 'ohcf': 6, 'dohcv': 7}
engine-location: {'front': 1, 'rear': 2}
drive-wheels: {'rwd': 1, 'fwd': 2, '4wd': 3}
body-style: {'convertible': 1, 'hatchback': 2, 'sedan': 3, 'wagon': 4, 'hardtop': 5}
aspiration: {'std': 1, 'turbo': 2}
fuel-type: {'gas': 1, 'diesel': 2}
make: {'alfa-romero': 1, 'audi': 2, 'bmw': 3, 'chevrolet': 4, 'dodge': 5, 'honda': 6, 'isuzu': 7, 'jaguar': 8, 'mazda': 9, 'mercedes
```

```
1 # checking the values of the dataframe
2 autom_df.head()
```

	price	highway-mpg	city-mpg	peak-rpm	horsepower	compression-ratio	stroke	bore	fuel-system	er
0	13495.0	27	21	5000.0	111.0	9.0	2.68	3.47	1	
1	16500.0	27	21	5000.0	111.0	9.0	2.68	3.47	1	
2	16500.0	26	19	5000.0	154.0	9.0	3.47	2.68	1	
3	13950.0	30	24	5500.0	102.0	10.0	3.40	3.19	1	
4	17450.0	22	18	5500.0	115.0	8.0	3.40	3.19	1	

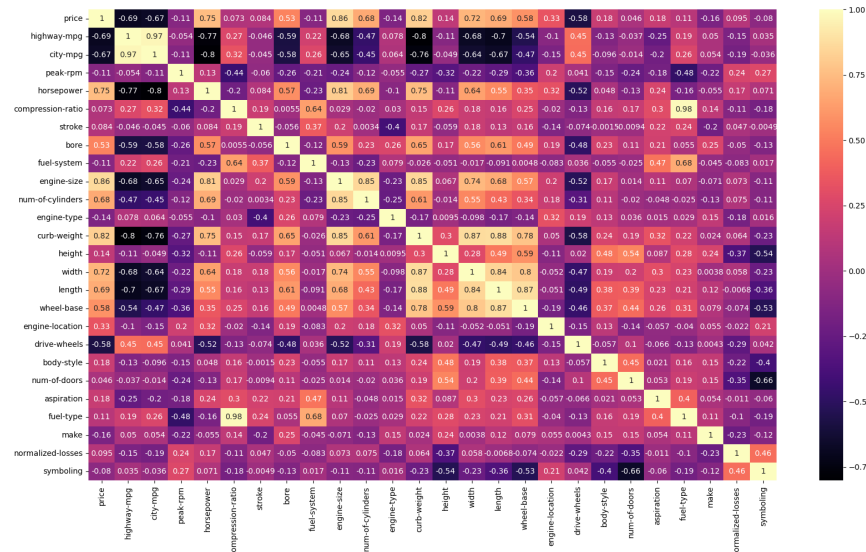
5 rows × 26 columns

```

1 # importing the necessary libraries for plotting
2 %matplotlib inline
3 import seaborn as sns
4 import matplotlib.pyplot as plt
5
6 # creating a heatmap to check the correlation of each categories between one another
7 plt.figure(figsize=(20, 11))
8 sns.heatmap(autom_df.corr(), annot=True, cmap='magma')

```

<Axes: >



Linear Regression Model

```

1 X = autom_df.drop('price', axis=1)
2 y = autom_df['price']

```

```
1 print("X", X.shape, "\ny=", y.shape)
```

```

X (205, 25)
y= (205,)

```

Train Test Splitting

```

1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
5 %matplotlib inline
6
7 from sklearn.model_selection import train_test_split
8 from sklearn import metrics
9 from sklearn.linear_model import LinearRegression

```

```
1 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=101)
```

```
1 X_train.shape

(143, 25)

1 X_test.shape

(62, 25)

1 model = LinearRegression()

1 model.fit(X_train, y_train)
```

▼ LinearRegression

LinearRegression()

▼ Model Evaluation

```
1 model.coef_

array([-1.50604087e+02,  7.37133859e+00,  1.16315603e+00, -2.93324387e+01,
        9.93470741e+02, -5.45481780e+03, -1.06091424e+04, -1.70268232e+02,
        2.05325227e+02, -3.00743877e+03, -5.28868529e+01,  1.41070253e+00,
        2.53044661e+02,  8.90488886e+02,  2.38868227e+01, -8.73652339e+01,
        1.54931808e+04, -2.22018286e+03, -4.10783975e+02,  3.27022203e+02,
        3.14166895e+03, -1.29759009e+04, -1.27546963e+02, -7.53386017e+00,
        3.34227339e+02])

1 pd.DataFrame(model.coef_, X.columns, columns=['Coedicients'])
```

	Coedicients	
highway-mpg	-150.604087	
city-mpg	7.371339	
peak-rpm	1.163156	
horsepower	-29.332439	
compression-ratio	993.470741	
stroke	-5454.817804	
bore	-10609.142408	
fuel-system	-170.268232	
engine-size	205.325227	
num-of-cylinders	-3007.438769	
engine-type	-52.886853	
curb-weight	1.410703	
height	253.044661	
width	890.488886	
length	23.886823	
wheel-base	-87.365234	
engine-location	15493.180825	
drive-wheels	-2220.182858	
body-style	-410.783975	
num-of-doors	327.022203	
aspiration	3141.668953	
fuel-type	-12975.900873	
make	-127.546963	
normalized-losses	-7.533860	
symboling	334.227339	

▼ Prediction from our Model

```
1 y_pred = model.predict(X_test)
```

Regression Evaluation Metrics

```
1 MAE = metrics.mean_absolute_error(y_test, y_pred)
2 MSE = metrics.mean_squared_error(y_test, y_pred)
3 RMSE = np.sqrt(MSE)
```

```
1 MAE
```

```
2554.5515500553283
```

```
1 MSE
```

```
14903114.988000777
```

```
1 RMSE
```

```
3860.4552824765083
```

```
1 autom_df['price'].mean()
```


```
13150.307317073171
```

Residual Histogram

```
1 test_residual = y_test - y_pred
```


```
1 pd.DataFrame({'Error Values': (test_residual)}).plot.kde()
```

<Axes: ylabel='Density'>




```
1 sns.displot(test_residual, bins=25, kde=True)
```


<seaborn.axisgrid.FacetGrid at 0x78a7dd329240>



```
1 sns.scatterplot(x=y_test, y=test_residual)
2 plt.axhline(y=0, color='r', ls='--')
```

<matplotlib.lines.Line2D at 0x78a7dd352200>



Wine

```
1 # Import the dataset into your code
2 from ucimlrepo import fetch_ucirepo
3
4 # fetch dataset
5 wine = fetch_ucirepo(id=109)
6
7 # data (as pandas dataframes)
8 X = wine.data.features
9 y = wine.data.targets
10
11 # metadata
12 print(wine.metadata)
13
14 # variable information
15 print(wine.variables)
```

```
{'uci_id': 109, 'name': 'Wine', 'repository_url': 'https://archive.ics.uci.edu/dataset/109/wine', 'data_url': 'https://archive.ics.uci.edu/dataset/109/wine'}
   name    role    type demographic \
0    class  Target  Categorical      None
1  Alcohol  Feature  Continuous      None
2  Malicacid Feature  Continuous      None
```

3	Ash	Feature	Continuous	None
4	Alcalinity_of_ash	Feature	Continuous	None
5	Magnesium	Feature	Integer	None
6	Total_phenols	Feature	Continuous	None
7	Flavanoids	Feature	Continuous	None
8	Nonflavanoid_phenols	Feature	Continuous	None
9	Proanthocyanins	Feature	Continuous	None
10	Color_intensity	Feature	Continuous	None
11	Hue	Feature	Continuous	None
12	0D280_0D315_of_diluted_wines	Feature	Continuous	None
13	Proline	Feature	Integer	None

	description	units	missing_values
0	None	None	no
1	None	None	no
2	None	None	no
3	None	None	no
4	None	None	no
5	None	None	no
6	None	None	no
7	None	None	no
8	None	None	no
9	None	None	no
10	None	None	no
11	None	None	no
12	None	None	no
13	None	None	no

```
1 # displaying the 'X' dataframe
2 X.head()
```

	Alcohol	Malicacid	Ash	Alcalinity_of_ash	Magnesium	Total_phenols	Flavanoids
0	14.23	1.71	2.43	15.6	127	2.80	3.06
1	13.20	1.78	2.14	11.2	100	2.65	2.76
2	13.16	2.36	2.67	18.6	101	2.80	3.24
3	14.37	1.95	2.50	16.8	113	3.85	3.49
4	13.24	2.59	2.87	21.0	118	2.80	2.69

Next steps:

View recommended plots

```
1 # displaying the 'y' dataframe
2 y.head()
```

	class
0	1
1	1
2	1
3	1
4	1

Next steps:

View recommended plots

```
1 # importing all the necessary libraries
2 import pandas as pd
3 import numpy as np
4
5 # concatenating all the 'X' and 'y' dataframe
6 wine_df = pd.concat([X, y], axis=1)
7 wine_df
```

	Alcohol	Malicacid	Ash	Alcalinity_of_ash	Magnesium	Total_phenols	Flavanoid:
0	14.23	1.71	2.43	15.6	127	2.80	3.06
1	13.20	1.78	2.14	11.2	100	2.65	2.76
2	13.16	2.36	2.67	18.6	101	2.80	3.24
3	14.37	1.95	2.50	16.8	113	3.85	3.46
4	13.24	2.59	2.87	21.0	118	2.80	2.69
...
173	13.71	5.65	2.45	20.5	95	1.68	0.67
174	13.40	3.91	2.48	23.0	102	1.80	0.71
175	13.27	4.28	2.26	20.0	120	1.59	0.69
176	13.17	2.59	2.37	20.0	120	1.65	0.68
177	14.13	4.10	2.74	24.5	96	2.05	0.76

178 rows × 14 columns

Next steps: [View recommended plots](#)

```
1 # checking if there are any duplicates
2 countDuplicate(wine_df)

'No Duplicates Found!'

1 # checking there are any nulls in the dataframe
2 wine_df.isnull().sum()

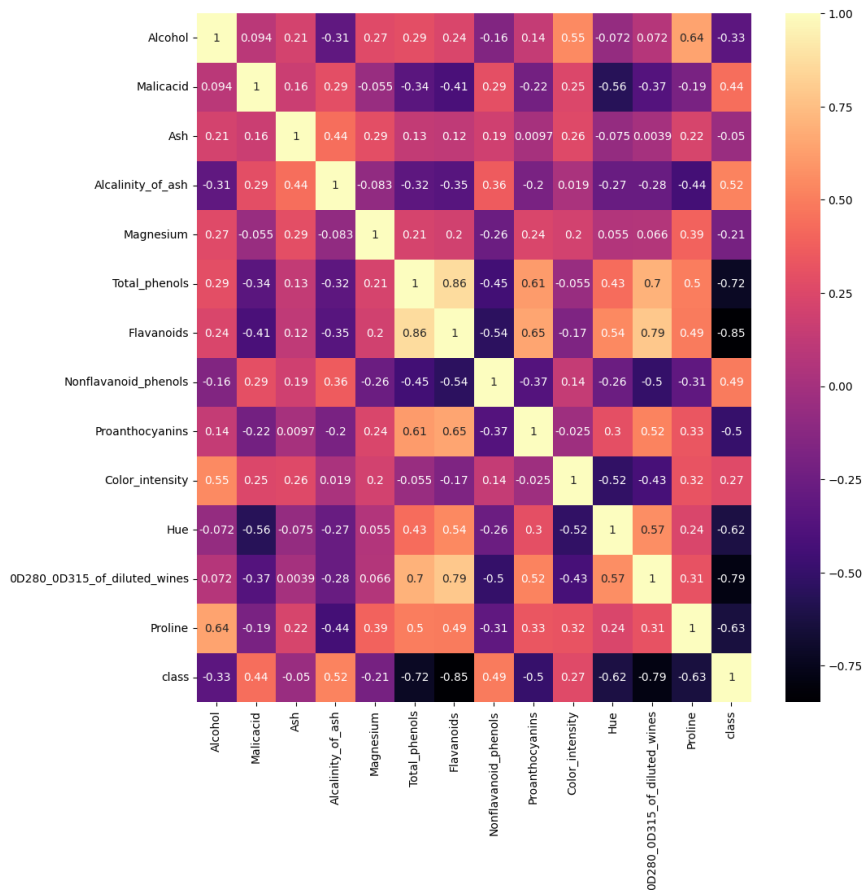
Alcohol      0
Malicacid    0
Ash          0
Alcalinity_of_ash  0
Magnesium    0
Total_phenols  0
Flavanoids   0
Nonflavanoid_phenols  0
Proanthocyanins  0
Color_intensity  0
Hue          0
0D280_0D315_of_diluted_wines  0
Proline      0
class       0
dtype: int64

1 # checking if there are any categorical values in the dataframe
2 wine_df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 178 entries, 0 to 177
Data columns (total 14 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   Alcohol                                   178 non-null    float64
1   Malicacid                                178 non-null    float64
2   Ash                                       178 non-null    float64
3   Alcalinity_of_ash                        178 non-null    float64
4   Magnesium                                178 non-null    int64
5   Total_phenols                            178 non-null    float64
6   Flavanoids                              178 non-null    float64
7   Nonflavanoid_phenols                    178 non-null    float64
8   Proanthocyanins                         178 non-null    float64
9   Color_intensity                         178 non-null    float64
10  Hue                                      178 non-null    float64
11  0D280_0D315_of_diluted_wines            178 non-null    float64
12  Proline                                  178 non-null    int64
13  class                                    178 non-null    int64
dtypes: float64(11), int64(3)
memory usage: 19.6 KB

1 plt.figure(figsize=(11,11))
2 sns.heatmap(wine_df.corr(), annot=True, cmap='magma')
```

<Axes: >



Logistic Regression

```

1 # importing libraries
2 import numpy as np # linear algebra
3 import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)
4 import matplotlib.pyplot as plt # data visualization
5 import seaborn as sns # statistical data visualization
6 %matplotlib inline

```

```

1 import warnings
2 warnings.filterwarnings('ignore')

```

```
1 print(round(wine_df.describe()),2)
```

	Alcohol	Malicacid	Ash	Alcalinity_of_ash	Magnesium	Total_phenols	\
count	178.0	178.0	178.0	178.0	178.0	178.0	
mean	13.0	2.0	2.0	19.0	100.0	2.0	
std	1.0	1.0	0.0	3.0	14.0	1.0	

min	11.0	1.0	1.0	11.0	70.0	1.0
25%	12.0	2.0	2.0	17.0	88.0	2.0
50%	13.0	2.0	2.0	20.0	98.0	2.0
75%	14.0	3.0	3.0	22.0	107.0	3.0
max	15.0	6.0	3.0	30.0	162.0	4.0

	Flavanoids	Nonflavanoid_phenols	Proanthocyanins	Color_intensity	\
count	178.0	178.0	178.0	178.0	178.0
mean	2.0	0.0	2.0	5.0	
std	1.0	0.0	1.0	2.0	
min	0.0	0.0	0.0	1.0	
25%	1.0	0.0	1.0	3.0	
50%	2.0	0.0	2.0	5.0	
75%	3.0	0.0	2.0	6.0	
max	5.0	1.0	4.0	13.0	

	Hue	0D280_0D315_of_diluted_wines	Proline	class
count	178.0	178.0	178.0	178.0
mean	1.0	3.0	747.0	2.0
std	0.0	1.0	315.0	1.0
min	0.0	1.0	278.0	1.0
25%	1.0	2.0	500.0	1.0
50%	1.0	3.0	674.0	2.0
75%	1.0	3.0	985.0	3.0
max	2.0	4.0	1680.0	3.0

```
1 wine_df['Alcalinity_of_ash'].min()

10.6
```

Malicacid:

- Upper bound: 4.5 and Max: 6.0, potential upper bound outlier

Alcalinity_of_ash:

- Upper bound: 29.5 and Max: 30, potential upper bound outlier

Magnesium:

- Upper bound: 135.5 and Max: 162, potential upper bound outlier

Proanthocyanins

- Upper bound: 3.5 and Max: 4, potential upper bound outlier

Color_intensity:

- Upper bound: 10.5 and Max: 13, potential upper bound outlier

Hue:

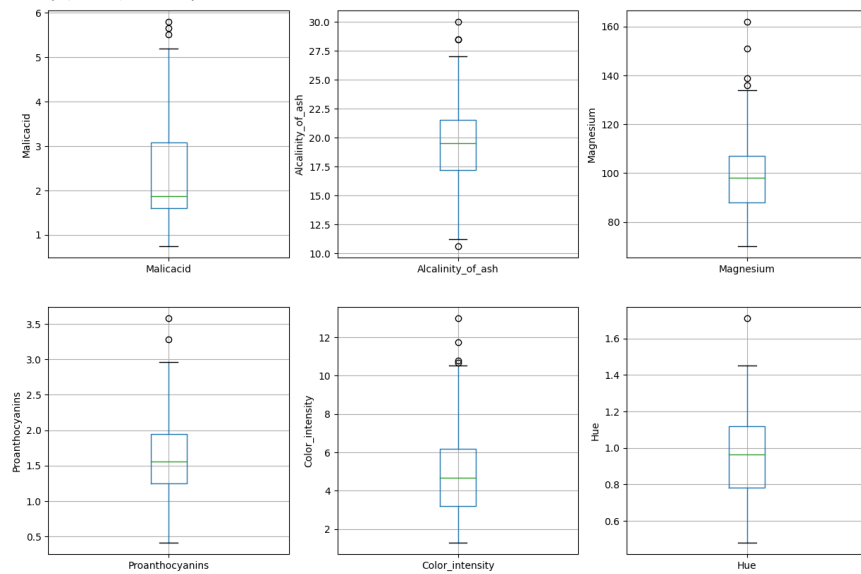
- Upper bound: 1 and Max: 2, potential upper bound outlier
- Lower bound: 1 and Min: 0, potential lower bound outlier

```

1 # draw boxplots to visualize the outliers
2 plt.figure(figsize=(15,10))
3
4 plt.subplot(2, 3, 1)
5 fig = wine_df.boxplot(column='Malicacid')
6 fig.set_title('')
7 fig.set_ylabel('Malicacid')
8
9 plt.subplot(2, 3, 2)
10 fig = wine_df.boxplot(column='Alcalinity_of_ash')
11 fig.set_title('')
12 fig.set_ylabel('Alcalinity_of_ash')
13
14 plt.subplot(2, 3, 3)
15 fig = wine_df.boxplot(column='Magnesium')
16 fig.set_title('')
17 fig.set_ylabel('Magnesium')
18
19 plt.subplot(2, 3, 4)
20 fig = wine_df.boxplot(column='Proanthocyanins')
21 fig.set_title('')
22 fig.set_ylabel('Proanthocyanins')
23
24 plt.subplot(2, 3, 5)
25 fig = wine_df.boxplot(column='Color_intensity')
26 fig.set_title('')
27 fig.set_ylabel('Color_intensity')
28
29 plt.subplot(2, 3, 6)
30 fig = wine_df.boxplot(column='Hue')
31 fig.set_title('')
32 fig.set_ylabel('Hue')

```

Text(0, 0.5, 'Hue')




```

1 # plot histogram to check distribution
2 # this is to determine which step would be next
3 plt.figure(figsize=(15,6))
4
5 plt.subplot(2, 3, 1)
6 fig = wine_df.Malicacid.hist(bins=10)
7 fig.set_xlabel('Malicacid')
8 fig.set_ylabel('')
9
10 plt.subplot(2, 3, 2)
11 fig = wine_df.Alcalinity_of_ash.hist(bins=10)
12 fig.set_xlabel('Alcalinity_of_ash')
13 fig.set_ylabel('')
14
15 plt.subplot(2, 3, 3)
16 fig = wine_df.Magnesium.hist(bins=10)
17 fig.set_xlabel('Magnesium')
18 fig.set_ylabel('')
19
20 plt.subplot(2, 3, 4)
21 fig = wine_df.Proanthocyanins.hist(bins=10)
22 fig.set_xlabel('Proanthocyanins')
23 fig.set_ylabel('')
24
25 plt.subplot(2, 3, 5)
26 fig = wine_df.Color_intensity.hist(bins=10)
27 fig.set_xlabel('Color_intensity')
28 fig.set_ylabel('')
29
30 plt.subplot(2, 3, 6)
31 fig = wine_df.Hue.hist(bins=10)
32 fig.set_xlabel('Hue')
33 fig.set_ylabel('')

```

Text(0, 0.5, '')



```

1 # since all 4 are skewed, next step would be interquartile range to find the outliers
2 IQR = wine_df['Malicacid'].quantile(0.75) - wine_df['Malicacid'].quantile(0.25)
3 Lower_fence = wine_df['Malicacid'].quantile(0.25) - (IQR * 1.5)
4 Upper_fence = wine_df['Malicacid'].quantile(0.75) + (IQR * 1.5)
5 print(f"Malicacid outliers are values < {Lower_fence} or > {Upper_fence}")

```

Malicacid outliers are values < -0.6174999999999997 or > 5.3025

```

1 IQR = wine_df['Alcalinity_of_ash'].quantile(0.75) - wine_df['Alcalinity_of_ash'].quantile(0.25)
2 Lower_fence = wine_df['Alcalinity_of_ash'].quantile(0.25) - (IQR * 1.5)
3 Upper_fence = wine_df['Alcalinity_of_ash'].quantile(0.75) + (IQR * 1.5)
4 print(f"Alcalinity_of_ash outliers are values < {Lower_fence} or > {Upper_fence}")

```

Alcalinity_of_ash outliers are values < 10.749999999999998 or > 27.950000000000003

```

1 IQR = wine_df['Magnesium'].quantile(0.75) - wine_df['Magnesium'].quantile(0.25)
2 Lower_fence = wine_df['Magnesium'].quantile(0.25) - (IQR * 1.5)
3 Upper_fence = wine_df['Magnesium'].quantile(0.75) + (IQR * 1.5)
4 print(f"Magnesium outliers are values < {Lower_fence} or > {Upper_fence}")

```

```

Magnesium outliers are values < 59.5 or > 135.5

```

```

1 IQR = wine_df['Proanthocyanins'].quantile(0.75) - wine_df['Proanthocyanins'].quantile(0.25)
2 Lower_fence = wine_df['Proanthocyanins'].quantile(0.25) - (IQR * 1.5)
3 Upper_fence = wine_df['Proanthocyanins'].quantile(0.75) + (IQR * 1.5)
4 print(f"Proanthocyanins outliers are values < {Lower_fence} or > {Upper_fence}")

```

```

Proanthocyanins outliers are values < 0.2000000000000018 or > 3.0

```

```

1 IQR = wine_df['Color_intensity'].quantile(0.75) - wine_df['Color_intensity'].quantile(0.25)
2 Lower_fence = wine_df['Color_intensity'].quantile(0.25) - (IQR * 1.5)
3 Upper_fence = wine_df['Color_intensity'].quantile(0.75) + (IQR * 1.5)
4 print(f"Color_intensity outliers are values < {Lower_fence} or > {Upper_fence}")

```

```

Color_intensity outliers are values < -1.250000000000009 or > 10.670000000000002

```

```

1 IQR = wine_df['Hue'].quantile(0.75) - wine_df['Hue'].quantile(0.25)
2 Lower_fence = wine_df['Hue'].quantile(0.25) - (IQR * 1.5)
3 Upper_fence = wine_df['Hue'].quantile(0.75) + (IQR * 1.5)
4 print(f"Hue outliers are values < {Lower_fence} or > {Upper_fence}")

```

```

Hue outliers are values < 0.276249999999998 or > 1.626250000000002

```

```

1 # Declare feature vector and target variable
2 X = wine_df.drop(['class'], axis=1)
3 y = wine_df['class']

```

```

1 # Split data into separate training and testing set
2 from sklearn.model_selection import train_test_split
3
4 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=0)

```

```

1 # check the shape of X_train and X_test
2 X_train.shape, X_test.shape

```

```

((142, 13), (36, 13))

```

```

1 X_train.dtypes

```

```

Alcohol          float64
Malicacid        float64
Ash              float64
Alcalinity_of_ash float64
Magnesium        int64
Total_phenols    float64
Flavanoids       float64
Nonflavanoid_phenols float64
Proanthocyanins  float64
Color_intensity  float64
Hue              float64
0D280_0D315_of_diluted_wines float64
Proline          int64
dtype: object

```

```

1 categorical = [i for i in X_train.columns if X_train[i].dtypes=='O']
2 categorical

```

```

[]

```

Since there are no categorical values in the given dataset we would proceed with checking the nulls

```

1 X_train.isnull().sum()

```

```

Alcohol          0
Malicacid        0
Ash              0
Alcalinity_of_ash 0
Magnesium        0
Total_phenols    0
Flavanoids       0
Nonflavanoid_phenols 0
Proanthocyanins  0
Color_intensity  0

```



```
Hue
0D280_0D315_of_diluted_wines
Proline
dtype: int64

1 X_test.isnull().any()

Alcohol
Malicacid
Ash
Alcalinity_of_ash
Magnesium
Total_phenols
Flavanoids
Nonflavanoid_phenols
Proanthocyanins
Color_intensity
Hue
0D280_0D315_of_diluted_wines
Proline
dtype: bool

1 def max_value(df3, variable, top):
2     return np.where(df3[variable]>top, top, df3[variable])
3
4 for df3 in [X_train, X_test]:
5     df3['Malicacid'] = max_value(df3, 'Malicacid', 5.30)
6     df3['Alcalinity_of_ash'] = max_value(df3, 'Alcalinity_of_ash', 27.95)
7     df3['Magnesium'] = max_value(df3, 'Magnesium', 135.5)
8     df3['Proanthocyanins'] = max_value(df3, 'Proanthocyanins', 3)
9     df3['Color_intensity'] = max_value(df3, 'Color_intensity', 10.67)
10    df3['Hue'] = max_value(df3, 'Hue', 1.63)

1 X_train['Malicacid'].max(), X_test['Malicacid'].max()

(5.3, 5.3)

1 X_train['Alcalinity_of_ash'].max(), X_test['Alcalinity_of_ash'].max()

(27.95, 27.95)

1 X_train['Magnesium'].max(), X_test['Magnesium'].max()

(135.5, 132.0)

1 X_train['Proanthocyanins'].max(), X_test['Proanthocyanins'].max()

(3.0, 2.45)

1 X_train['Color_intensity'].max(), X_test['Color_intensity'].max()

(10.67, 10.67)

1 X_train['Hue'].max(), X_test['Hue'].max()

(1.63, 1.38)

1 X_train.describe()
```

	Alcohol	Malicacid	Ash	Alcalinity_of_ash	Magnesium	Total_phenols	Flavanoids	Nonflavanoid_phenols	Proanthocyanins
count	142.000000	142.000000	142.000000	142.000000	142.000000	142.000000	142.000000	142.000000	142.000000
mean	12.984859	2.368662	2.366901	19.536620	99.739437	2.258662	1.949155	0.363521	1.600000
std	0.807175	1.104345	0.269684	3.392529	13.154391	0.611691	0.975921	0.127709	0.570000
min	11.030000	0.740000	1.360000	10.600000	70.000000	1.100000	0.470000	0.130000	0.420000
25%	12.347500	1.602500	2.222500	17.250000	89.000000	1.705000	1.037500	0.270000	1.240000
50%	13.040000	1.895000	2.360000	19.500000	98.000000	2.210000	2.035000	0.340000	1.550000
75%	13.637500	3.222500	2.560000	21.500000	106.750000	2.735000	2.760000	0.450000	1.950000
max	14.750000	5.300000	3.220000	27.950000	135.500000	3.880000	3.740000	0.660000	3.000000

```
1 # Feature scaling
2 X_train.describe()
```

	Alcohol	Malicacid	Ash	Alcalinity_of_ash	Magnesium	Total_phenols	Flavanoids	Nonflavanoid_phenols	Proanthocya
count	142.000000	142.000000	142.000000	142.000000	142.000000	142.000000	142.000000	142.000000	142.00
mean	12.984859	2.368662	2.366901	19.536620	99.739437	2.258662	1.949155	0.363521	1.60
std	0.807175	1.104345	0.269684	3.392529	13.154391	0.611691	0.975921	0.127709	0.57
min	11.030000	0.740000	1.360000	10.600000	70.000000	1.100000	0.470000	0.130000	0.42
25%	12.347500	1.602500	2.222500	17.250000	89.000000	1.705000	1.037500	0.270000	1.24
50%	13.040000	1.895000	2.360000	19.500000	98.000000	2.210000	2.035000	0.340000	1.55
75%	13.637500	3.222500	2.560000	21.500000	106.750000	2.735000	2.760000	0.450000	1.95
max	14.750000	5.300000	3.220000	27.950000	135.500000	3.880000	3.740000	0.660000	3.00

```
1 cols = X_train.columns

1 from sklearn.preprocessing import MinMaxScaler
2
3 scaler = MinMaxScaler()
4 X_train = scaler.fit_transform(X_train)
5 X_test = scaler.transform(X_test)

1 X_train = pd.DataFrame(X_train, columns=[cols])

1 X_test = pd.DataFrame(X_test, columns=[cols])

1 X_train.describe()
```

	Alcohol	Malicacid	Ash	Alcalinity_of_ash	Magnesium	Total_phenols	Flavanoids	Nonflavanoid_phenols	Proanthocya
count	142.000000	142.000000	142.000000	142.000000	142.000000	142.000000	142.000000	142.000000	142.00
mean	0.525500	0.357163	0.541345	0.515079	0.454037	0.416785	0.452341	0.440606	0.45
std	0.216983	0.242181	0.144991	0.195535	0.200830	0.220033	0.298447	0.240960	0.22
min	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	0.00
25%	0.354167	0.189145	0.463710	0.383285	0.290076	0.217626	0.173547	0.264151	0.31
50%	0.540323	0.253289	0.537634	0.512968	0.427481	0.399281	0.478593	0.396226	0.43
75%	0.700941	0.544408	0.645161	0.628242	0.561069	0.588129	0.700306	0.603774	0.55
max	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.000000	1.00

```
1 # Model training
2 # train a logistic regression model on the training set
3 from sklearn.linear_model import LogisticRegression
4
5 # instantiate the model
6 logreg = LogisticRegression(solver='liblinear', random_state=0)
7
8 #fit the model
9 logreg.fit(X_train, y_train)
```

▼

LogisticRegression

LogisticRegression(random_state=0, solver='liblinear')

```
1 # Predicting results
2 y_pred_test = logreg.predict(X_test)
3 y_pred_test

array([1, 3, 2, 1, 2, 2, 1, 3, 2, 2, 3, 3, 1, 2, 3, 2, 1, 1, 3, 1, 2, 1,
       1, 2, 2, 2, 2, 2, 2, 3, 1, 1, 2, 1, 1, 1])

1 # predict proba: predicts possibilities for the target variable
2 logreg.predict_proba(X_test)[:,:0]

array([0.84608661, 0.08048314, 0.33554459, 0.80152316, 0.22039744,
       0.24522358, 0.87468447, 0.03848718, 0.15948773, 0.05898538,
```

```
0.14507607, 0.0358635 , 0.93262766, 0.47681744, 0.07875799,
0.12976804, 0.79007758, 0.95638587, 0.08820506, 0.84199176,
0.47786966, 0.67550962, 0.47265773, 0.23584062, 0.08901126,
0.15152147, 0.20646317, 0.04988917, 0.07608428, 0.07232845,
0.83715721, 0.86643051, 0.0791357 , 0.81872064, 0.87889396,
0.67623076])
```

```
1 logreg.predict_proba(X_test)[: ,1]
```

```
array([0.10773965, 0.04382495, 0.65481533, 0.1525236 , 0.66114073,
0.74433276, 0.07438344, 0.12116596, 0.78373262, 0.80051396,
0.1289791 , 0.07885248, 0.03352518, 0.51535383, 0.06169256,
0.85700534, 0.14896474, 0.01820935, 0.40806803, 0.14198268,
0.51407897, 0.25405006, 0.43617131, 0.73190936, 0.59297914,
0.78249892, 0.7421782 , 0.86155799, 0.73512681, 0.0449725 ,
0.12680024, 0.09818553, 0.64410061, 0.05781283, 0.08848967,
0.30025453])
```

```
1 logreg.predict_proba(X_test)[: ,2]
```

```
array([0.04617374, 0.87569191, 0.00964009, 0.04595323, 0.11846182,
0.01044366, 0.05093209, 0.84034686, 0.05677965, 0.14050066,
0.72594482, 0.88528401, 0.03384715, 0.00782873, 0.85954945,
0.01322662, 0.06095768, 0.02540478, 0.50372691, 0.01602556,
0.00805136, 0.07044032, 0.09117095, 0.03225002, 0.31800959,
0.06597962, 0.05135863, 0.08855284, 0.18878891, 0.88269905,
0.03604255, 0.03538396, 0.27676368, 0.12346652, 0.03261637,
0.02351471])
```

```
1 # Check accuracy score
2 from sklearn.metrics import accuracy_score
3
4 print('Model accuracy score: {0:0.4f}'.format(accuracy_score(y_test, y_pred_test)))
```

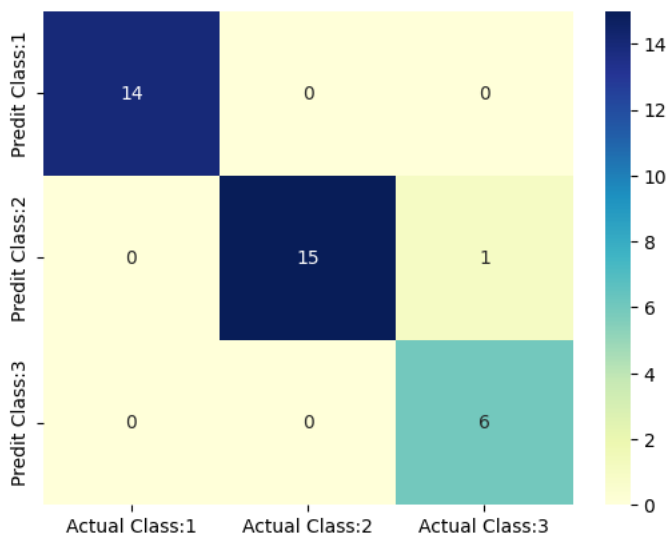
```
Model accuracy score: 0.9722
```

```
1 # check for overfitting and underfitting
2 # print the scores on training and test set
3 train_score = logreg.score(X_train, y_train)
4 test_score = logreg.score(X_test, y_test)
5 print(f'Train set score: {train_score}')
6 print(f'Test set score: {test_score}')
```

```
Train set score: 0.9788732394366197
Test set score: 0.9722222222222222
```

```
1 # visualize confusion matrix with seaborn heatmap
2 cm_matrix = pd.DataFrame(data=cm, columns=['Actual Class:1', 'Actual Class:2', 'Actual Class:3'],
3                             index=['Predit Class:1', 'Predit Class:2', 'Predit Class:3'])
4 sns.heatmap(cm_matrix, annot=True, fmt='d', cmap='YlGnBu')
```

<Axes: >



```
1 # classification metrics
2 from sklearn.metrics import classification_report
3 print(classification_report(y_test, y_pred_test))
```

```
precision    recall  f1-score   support
```

1	1.00	1.00	1.00	14
2	1.00	0.94	0.97	16
3	0.86	1.00	0.92	6
accuracy			0.97	36
macro avg	0.95	0.98	0.96	36
weighted avg	0.98	0.97	0.97	36

```
1 y_pred_prob = logreg.predict_proba(X_test)[0:10]
2 y_pred_prob
```

```
array([[0.84608661, 0.10773965, 0.04617374],
       [0.08048314, 0.04382495, 0.87569191],
       [0.33554459, 0.65481533, 0.00964009],
       [0.80152316, 0.1525236 , 0.04595323],
       [0.22039744, 0.66114073, 0.11846182],
       [0.24522358, 0.74433276, 0.01044366],
       [0.87468447, 0.07438344, 0.05093209],
       [0.03848718, 0.12116596, 0.84034686],
       [0.15948773, 0.78373262, 0.05677965],
       [0.05898538, 0.80051396, 0.14050066]])
```


```
1 # print the first 10 predicted probabilities
2 logreg.predict_proba(X_test)[0:10, 1]
```

```
array([0.10773965, 0.04382495, 0.65481533, 0.1525236 , 0.66114073,
       0.74433276, 0.07438344, 0.12116596, 0.78373262, 0.80051396])
```

```
1 # store the predicted probabilities
2 y_pred1 = logreg.predict_proba(X_test)[: ,1]
```

```
1 # plot histogram of predicted probabilities
2 # adjust the font size
3 plt.rcParams['font.size'] = 12
4
5 # plot histogram with 10 bins
6 plt.hist(y_pred1, bins=10)
7
8 # set the title of predicted probabilities
9 plt.title('Histogram of predicted probabilities of class')
10
11 # set the x-axis limit
12 plt.xlim(0,1)
13
14 plt.xlabel('Predicted probabilities of class')
15 plt.ylabel('Frequency')
```

```
Text(0, 0.5, 'Frequency')
```



```
1 # k-Fold Cross Validation
2 from sklearn.model_selection import cross_val_score
3
4 scores = cross_val_score(logreg, X_train, y_train, cv = 5, scoring='accuracy')
5 print('Cross-validation scores:{}'.format(scores))
```

```

Cross-validation scores:[0.93103448 0.96551724 0.96428571 1.          0.96428571]

1 # compute Average cross-validation score
2 score_mean = scores.mean()
3 print(f'Average cross-validation score: {score_mean}')

Average cross-validation score: 0.9650246305418719

1 # Hyperparameter Optimization using GridSearchCV
2 from sklearn.model_selection import GridSearchCV
3
4 parameters = [{'penalty': ['l1', 'l2']],
5               {'C':[1, 10, 100, 1000]}]
6
7 grid_search = GridSearchCV(estimator = logreg,
8                             param_grid = parameters,
9                             scoring = 'accuracy',
10                            cv = 5,
11                            verbose=0)
12
13 grid_search.fit(X_train, y_train)

```

```

GridSearchCV
└─ estimator: LogisticRegression
   └─ LogisticRegression
      LogisticRegression(random_state=0, solver='liblinear')

```

```

1 # examine the best model
2 # best score achieved during the GridSearchCV
3 print('GridSearch CV best score : {:.4f}\n\n'.format(grid_search.best_score_))
4
5 # print parameters that give the best results
6 print('Parameters that give the best results :', '\n\n', (grid_search.best_params_))
7
8 # print estimator that was chosen by the GridSearch
9 print('\n\nEstimator that was chosen by the search :', '\n\n', (grid_search.best_estimator_))

```

GridSearch CV best score : 0.9650

Parameters that give the best results :

```
{'C': 1}
```

Estimator that was chosen by the search :

```
LogisticRegression(C=1, random_state=0, solver='liblinear')
```

```

1 # calculate Gridsearch CV score on test set
2 print('Gridsearch CV score on test set: {:.4f}'.format(grid_search.score(X_test, y_test)))

```

Gridsearch CV score on test set: 0.9722

✓ Results and Conclusion

- The logistic regression model accuracy score is 0.9722. Since the value the train sets are above 0.9, we can conclude that the training is accurate on its predictions. Therefore, the model did an excellent work in predicting the class of each columns.

1