

✓ Hands-on Activity 8.1: Aggregating Data with Pandas

Name: Cuadra, Audrick Zander G.

Section: CPE22S3

Date: March 29, 2024

Submitted to: Engr. Roman Richard

About the data

After this activity, the student should be able to:

- Demonstrate querying and merging of dataframes
- Perform advanced calculations on dataframes
- Aggregate dataframes with pandas and numpy
- Work with time series data

8.1.2 Resources

- Computing Environment using Python 3.x
- Attached Datasets (under Instructional Materials)

8.1.3 Procedures

The procedures can be found in the canvas module. Check the following under topics:

- 8.1 Weather Data Collection
- 8.2 Querying and Merging
- 8.3 Dataframe Operations
- 8.4 Aggregations
- 8.5 Time Series

8.1.4 Data Analysis

The modules demonstrates methods that make use of several Python packages for data analysis. This module includes a number of libraries, including data cleaning, modification, and visualization. Additionally, by utilizing real-world data, it illustrates the significance of data analysis from the real world. Overall, the modules offer learners more than enough opportunities to learn about various data manipulation and analysis approaches.

8.1.5 Supplementary Activity

Using the CSV files provided and what we have learned so far in this module complete the following exercises:

1. With the earthquakes.csv file, select all the earthquakes in Japan with a magType of mb and a magnitude of 4.9 or greater.

```

1 # Used for the few first column as well as the structure of the dataframe
2 import pandas as p
3
4 earthquakes = p.read_csv('/content/earthquakes.csv')
5 earthquakes.head()

```

	mag	magType	time	place	tsunami	parsed_place
0	1.35	ml	1539475168010	9km NE of Aguanga, CA	0	California
1	1.29	ml	1539475129610	9km NE of Aguanga, CA	0	California
2	3.42	ml	1539475062610	8km NE of Aguanga, CA	0	California
3	0.44	ml	1539474978070	9km NE of Aguanga, CA	0	California
4	2.16	md	1539474716050	10km NW of Avenal, CA	0	California

```

1 # filters out the column name and column value type as well as mag limiter
2 earthquakes_data = earthquakes.query('magType == "mb" and mag >= 4.9')
3 earthquakes_data

```

	mag	magType	time	place	tsunami	parsed_place
227	5.2	mb	1539389603790	15km WSW of Pisco, Peru	0	Peru
229	4.9	mb	1539389546300	193km N of Qulansiyah, Yemen	0	Yemen
248	4.9	mb	1539382925190	151km S of Severo- Kuril'sk, Russia	0	Russia
258	5.1	mb	1539380306940	236km NNW of Kuril'sk, Russia	0	Russia
391	5.1	mb	1539337221080	Pacific-Antarctic Ridge	0	Pacific-Antarctic Ridge
...
9154	4.9	mb	1537268270010	Southwest Indian Ridge	0	Southwest Indian Ridge
9175	5.2	mb	1537262729590	126km N of Dili, East Timor	1	East Timor

2. Create bins for each full number of magnitude (for example, the first bin is 0-1, the second is 1-2, and so on) with a magType of ml and count how many are in each bin.

```

1 # filters out the the rows that contains the magType ml
2 earthquakes_ml = earthquakes.query('magType == "ml"')
3 earthquakes_ml

```

	mag	magType	time	place	tsunami	parsed_place
0	1.35	ml	1539475168010	9km NE of Aguanga, CA	0	California
1	1.29	ml	1539475129610	9km NE of Aguanga, CA	0	California
2	3.42	ml	1539475062610	8km NE of Aguanga, CA	0	California
3	0.44	ml	1539474978070	9km NE of Aguanga, CA	0	California
6	1.70	ml	1539473176017	105km W of Talkeetna, Alaska	0	Alaska
...
9325	0.51	ml	1537230344890	4km WNW of Julian, CA	0	California
9326	1.82	ml	1537230230260	4km W of Julian, CA	0	California
9328	1.00	ml	1537230135130	3km W of Julian, CA	0	California
9330	1.10	ml	1537229545350	9km NE of Aguanga, CA	0	California
9331	0.66	ml	1537228864470	9km NE of Aguanga, CA	0	California

6803 rows × 6 columns

```

1 # checks for the maximum value of mag to know how many bins should be made
2 max(earthquakes_ml.mag)

```

5.1

```

1 # creation of bins based on the specified conditions and displays value count
2 magnitude_binned = p.cut(
3     earthquakes_ml.mag, bins=6, labels=['0-1', '1-2', '2-3', '3-4', '4-5', '5-6']
4 )
5 magnitude_binned.value_counts()

```

```

2-3    3436
1-2    1889
3-4    1027
0-1     288
4-5     160
5-6       3

```

Name: mag, dtype: int64

3. Using the faang.csv file, group by the ticker and resample to monthly frequency. Make the following aggregations:

- Mean of the opening price
- Maximum of the high price
- Minimum of the low price

- Mean of the closing price
- Sum of the volume traded

```
1 # reads the csv file and checks the structure of the dataframe
2 faang = p.read_csv('/content/faang.csv', index_col= 'date', parse_dates=True)
3 faang.head()
```

	ticker	open	high	low	close	volume
date						
2018-01-02	FB	177.68	181.58	177.5500	181.42	18151903
2018-01-03	FB	181.88	184.78	181.3300	184.67	16886563
2018-01-04	FB	184.90	186.21	184.0996	184.33	13880896
2018-01-05	FB	185.59	186.90	184.9300	186.85	13574535
2018-01-08	FB	187.20	188.90	186.3300	188.28	17994726

```
1 # displayed value based on the what was asked
2 import numpy as ny
3
4 faang_monthly = faang.groupby('ticker').resample('M').agg({
5     'open': ny.mean,
6     'high': ny.max,
7     'low': ny.min,
8     'close': ny.mean,
9     'volume': ny.sum
10 })
11 faang_monthly
```

		open	high	low	close	volume
ticker	date					
AAPL	2018-01-31	170.714690	176.6782	161.5708	170.699271	659679440
	2018-02-28	164.562753	177.9059	147.9865	164.921884	927894473
	2018-03-31	172.421381	180.7477	162.4660	171.878919	713727447
	2018-04-30	167.332895	176.2526	158.2207	167.286924	666360147
	2018-05-31	182.635582	187.9311	162.7911	183.207418	620976206
	2018-06-30	186.605843	192.0247	178.7056	186.508652	527624365
	2018-07-31	188.065786	193.7650	181.3655	188.179724	393843881
	2018-08-31	210.460287	227.1001	195.0999	211.477743	700318837
	2018-09-30	220.611742	227.8939	213.6351	220.356353	678972040
	2018-10-31	219.489426	231.6645	204.4963	219.137822	789748068
	2018-11-30	190.828681	220.6405	169.5328	190.246652	961321947
	2018-12-31	164.537405	184.1501	145.9639	163.564732	898917007
AMZN	2018-01-31	1301.377143	1472.5800	1170.5100	1309.010952	96371290
	2018-02-28	1447.112632	1528.7000	1265.9300	1442.363158	137784020
	2018-03-31	1542.160476	1617.5400	1365.2000	1540.367619	130400151
	2018-04-30	1475.841905	1638.1000	1352.8800	1468.220476	129945743
	2018-05-31	1590.474545	1635.0000	1546.0200	1594.903636	71615299
	2018-06-30	1699.088571	1763.1000	1635.0900	1698.823810	85941510
	2018-07-31	1786.305714	1880.0500	1678.0600	1784.649048	97629820
	2018-08-31	1891.957826	2025.5700	1776.0200	1897.851304	96575676
	2018-09-30	1969.239474	2050.5000	1865.0000	1966.077895	94445693
	2018-10-31	1799.630870	2033.1900	1476.3600	1782.058261	183228552
	2018-11-30	1622.323810	1784.0000	1420.0000	1625.483810	139290208
	2018-12-31	1572.922105	1778.3400	1307.0000	1559.443158	154812304
FB	2018-01-31	184.364762	190.6600	175.8000	184.962857	495655736
	2018-02-28	180.721579	195.3200	167.1800	180.269474	516621991
	2018-03-31	173.449524	186.1000	149.0200	173.489524	996232472
	2018-04-30	164.163557	177.1000	150.5100	163.810476	751130388
	2018-05-31	181.910509	192.7200	170.2300	182.930000	401144183
	2018-06-30	194.974067	203.5500	186.4300	195.267619	387265765

4. Build a crosstab with the earthquake data between the tsunami column and the magType column. Rather than showing the frequency count, show the maximum magnitude that was observed for each combination. Put the magType along the columns.

```

2010-11-03    171.702007    107.1000    120.0000    171.000717    310100410
1 # made a list of values of each columns
2 # displayed the values of max magnitude for each combination
3 earthquakes_crosstab = p.crosstab(
4     index=earthquakes['tsunami'],
5     columns=earthquakes['magType'],
6     values=earthquakes['mag'],
7     aggfunc=ny.max
8 )
9 earthquakes_crosstab

```

magType	mb	mb_lg	md	mh	ml	ms_20	mw	mwb	mwr	mww
tsunami										
0	5.6	3.5	4.11	1.1	4.2	NaN	3.83	5.8	4.8	6.0
1	6.1	NaN	NaN	NaN	5.1	5.7	4.41	NaN	NaN	7.5

5. Calculate the rolling 60-day aggregations of OHLC data by ticker for the FAANG data. Use the same aggregations as exercise no. 3.

```

NFLX    2010-01-01    201.209200    200.0100    190.4200    202.900050    200377000
1 # checks the structure of the dataframe
2 faang

```

	ticker	open	high	low	close	volume
date						
2018-01-02	FB	177.68	181.58	177.5500	181.42	18151903
2018-01-03	FB	181.88	184.78	181.3300	184.67	16886563
2018-01-04	FB	184.90	186.21	184.0996	184.33	13880896
2018-01-05	FB	185.59	186.90	184.9300	186.85	13574535
2018-01-08	FB	187.20	188.90	186.3300	188.28	17994726
...
2018-12-24	GOOG	973.90	1003.54	970.1100	976.22	1590328
2018-12-26	GOOG	989.01	1040.00	983.0000	1039.46	2373270
2018-12-27	GOOG	1017.15	1043.89	997.0000	1043.88	2109777
2018-12-28	GOOG	1049.62	1055.56	1033.1000	1037.08	1413772
2018-12-31	GOOG	1050.96	1052.70	1023.5900	1035.61	1493722

1255 rows × 6 columns

```

1 # groups the OHLC aggregation by its ticker
2 faang.groupby('ticker').rolling('60D').agg({
3     'open': ny.mean,
4     'high': ny.max,
5     'low': ny.min,
6     'close': ny.mean,
7     'volume': ny.sum
8 })

```

		open	high	low	close	volume
ticker	date					
AAPL	2018-01-02	166.927100	169.0264	166.0442	168.987200	25555934.0
	2018-01-03	168.089600	171.2337	166.0442	168.972500	55073833.0
	2018-01-04	168.480367	171.2337	166.0442	169.229200	77508430.0
	2018-01-05	168.896475	172.0381	166.0442	169.840675	101168448.0
	2018-01-08	169.324680	172.2736	166.0442	170.080040	121736214.0
...
NFLX	2018-12-24	283.509250	332.0499	233.6800	281.931750	525657894.0
	2018-12-26	281.844500	332.0499	231.2300	280.777750	520444588.0
	2018-12-27	281.070488	332.0499	231.2300	280.162805	532679805.0
	2018-12-28	279.916341	332.0499	231.2300	279.461341	521968250.0
	2018-12-31	278.430769	332.0499	231.2300	277.451410	476309676.0

1255 rows × 5 columns

6. Create a pivot table of the FAANG data that compares the stocks. Put the ticker in the rows and show the averages of the OHLC and volume traded data.

```

1 # resets the index of the dataframe, transforming it back to default
2 # creates a pivot table that has an index ticker and gets the average of the specified
3 faang_pivot = faang.reset_index().pivot_table(
4     index='ticker',
5     values=['open', 'high', 'low', 'close'],
6     aggfunc='mean'
7 )
8
9 faang_pivot

```

	close	high	low	open
ticker				
AAPL	186.986218	188.906858	185.135729	187.038674
AMZN	1641.726175	1662.839801	1619.840398	1644.072669
FB	171.510936	173.615298	169.303110	171.454424
GOOG	1113.225139	1125.777649	1101.001594	1113.554104
NFLX	319.290299	325.224583	313.187273	319.620533

Next steps: ☒ View recommended plots

7. Calculate the Z-scores for each numeric column of Netflix's data (ticker is NFLX) using `apply()`.

```
1 # filters out all of the values with the ticker "NFLX"
2 faang_data = faang.query('ticker == "NFLX"')
3 faang_data
```

	ticker	open	high	low	close	volume
date						
2018-01-02	NFLX	196.10	201.6500	195.4200	201.070	10966889
2018-01-03	NFLX	202.05	206.2100	201.5000	205.050	8591369
2018-01-04	NFLX	206.20	207.0500	204.0006	205.630	6029616
2018-01-05	NFLX	207.25	210.0200	205.5900	209.990	7033240
2018-01-08	NFLX	210.02	212.5000	208.4400	212.050	5580178
...
2018-12-24	NFLX	242.00	250.6500	233.6800	233.880	9547616
2018-12-26	NFLX	233.92	254.5000	231.2300	253.670	14402735
2018-12-27	NFLX	250.11	255.5900	240.1000	255.565	12235217
2018-12-28	NFLX	257.94	261.9144	249.8000	256.080	10987286
2018-12-31	NFLX	260.16	270.1001	260.0000	267.660	13508920

251 rows × 6 columns

```
1 # checks the columns that have numerical datatype
2 faang.dtypes
```

```
ticker    object
open      float64
high      float64
low       float64
close     float64
volume    int64
dtype: object
```

```
1 # calls upon all of the numerical columns and performs z score operation to all of the
2 faang_data_z_score = faang_data.loc[
3     '2018', ['open', 'high', 'low', 'close', 'volume']
4 ].apply(
5     lambda x: x.sub(x.mean()).div(x.std())
6 )
7 faang_data_z_score
```

	open	high	low	close	volume
date					
2018-01-02	-2.500753	-2.516023	-2.410226	-2.416644	-0.088760
2018-01-03	-2.380291	-2.423180	-2.285793	-2.335286	-0.507606
2018-01-04	-2.296272	-2.406077	-2.234616	-2.323429	-0.959287
2018-01-05	-2.275014	-2.345607	-2.202087	-2.234303	-0.782331
2018-01-08	-2.218934	-2.295113	-2.143759	-2.192192	-1.038531
...
2018-12-24	-1.571478	-1.518366	-1.627197	-1.745946	-0.339003
2018-12-26	-1.735063	-1.439978	-1.677339	-1.341402	0.517040
2018-12-27	-1.407286	-1.417785	-1.495805	-1.302664	0.134868
2018-12-28	-1.248762	-1.289018	-1.297285	-1.292137	-0.085164
2018-12-31	-1.203817	-1.122354	-1.088531	-1.055420	0.359444

251 rows × 5 columns

8. Add event descriptions:

a.) Create a dataframe with the following three columns: ticker, date, and event. The columns should have the following values:

- ticker: 'FB'
- date: ['2018-07-25', '2018-03-19', '2018-03-20']
- event: ['Disappointing user growth announced after close.', 'Cambridge Analytica story', 'FTC investigation']

b.) Set the index to ['date', 'ticker']

c.) Merge this data with the FAANG data using an outer join

```
1 # creates a new dataframe for faang.csv
2 faang2 = p.read_csv(
3     '/content/faang.csv'
4 )
5 faang2
```

	ticker	date	open	high	low	close	volume
0	FB	2018-01-02	177.68	181.58	177.5500	181.42	18151903
1	FB	2018-01-03	181.88	184.78	181.3300	184.67	16886563
2	FB	2018-01-04	184.90	186.21	184.0996	184.33	13880896
3	FB	2018-01-05	185.59	186.90	184.9300	186.85	13574535
4	FB	2018-01-08	187.20	188.90	186.3300	188.28	17994726
...
1250	GOOG	2018-12-24	973.90	1003.54	970.1100	976.22	1590328
1251	GOOG	2018-12-26	989.01	1040.00	983.0000	1039.46	2373270
1252	GOOG	2018-12-27	1017.15	1043.89	997.0000	1043.88	2109777
1253	GOOG	2018-12-28	1049.62	1055.56	1033.1000	1037.08	1413772
1254	GOOG	2018-12-31	1050.96	1052.70	1023.5900	1035.61	1493722

1255 rows × 7 columns

```
1 # filters out the ticker "FB"
2 faang_fb = faang2.query('ticker == "FB"')
3 faang_fb
```

	ticker	date	open	high	low	close	volume
0	FB	2018-01-02	177.68	181.58	177.5500	181.42	18151903
1	FB	2018-01-03	181.88	184.78	181.3300	184.67	16886563
2	FB	2018-01-04	184.90	186.21	184.0996	184.33	13880896
3	FB	2018-01-05	185.59	186.90	184.9300	186.85	13574535
4	FB	2018-01-08	187.20	188.90	186.3300	188.28	17994726
...
246	FB	2018-12-24	123.10	129.74	123.0200	124.06	22066002
247	FB	2018-12-26	126.00	134.24	125.8900	134.18	39723370
248	FB	2018-12-27	132.44	134.99	129.6700	134.52	31202509
249	FB	2018-12-28	135.34	135.92	132.2000	133.20	22627569
250	FB	2018-12-31	134.45	134.64	129.9500	131.09	24625308

251 rows × 7 columns

```
1 # creates a new dataframe with the extracted date and ticker column
2 faang_new = faang_fb.filter(['date', 'ticker'])
3 faang_new
```

	date	ticker
0	2018-01-02	FB
1	2018-01-03	FB
2	2018-01-04	FB
3	2018-01-05	FB
4	2018-01-08	FB
...
246	2018-12-24	FB
247	2018-12-26	FB
248	2018-12-27	FB
249	2018-12-28	FB
250	2018-12-31	FB

251 rows × 2 columns

```

1 # writes a condition wherein the events would appear with a specific date
2 # otherwise the date's event would be "NaN"
3 date_new = faang_new['date'].isin(['2018-07-25', '2018-03-19', '2018-03-20'])
4 events = ['Disappointing user growth announced after close.',
5           'Cambridge Analytica story',
6           'FTC investigation']
7
8 faang_new['events'] = 'NaN'
9 faang_new.loc[date_new, 'events'] = ', '.join(events)
10
11 faang_new.query('date == ["2018-07-25", "2018-03-19", "2018-03-20"]')
```

	date	ticker	events
52	2018-03-19	FB	Disappointing user growth announced after clos...
53	2018-03-20	FB	Disappointing user growth announced after clos...
141	2018-07-25	FB	Disappointing user growth announced after clos...

9. Use the `transform()` method on the FAANG data to represent all the values in terms of the first date in the data. To do so, divide all the values for each ticker by the values for the first date in the data for that ticker. This is referred to as an index, and the data for the first date is the base (<https://ec.europa.eu/eurostat/statistics-explained/index.php/Beggins:Statisticalconcept-Indexandbaseyear>). When data is in this format, we can easily see growth over time. Hint: `transform()` can take a function name.

```
1 faang_new_index = faang.groupby('ticker').transform(  
2     lambda x: x.iloc[0]  
3 )  
4 faang_new_index
```

	open	high	low	close	volume
date					
2018-01-02	177.68	181.58	177.55	181.42	18151903.0
2018-01-03	177.68	181.58	177.55	181.42	18151903.0
2018-01-04	177.68	181.58	177.55	181.42	18151903.0
2018-01-05	177.68	181.58	177.55	181.42	18151903.0
2018-01-08	177.68	181.58	177.55	181.42	18151903.0
...
2018-12-24	1048.34	1066.94	1045.23	1065.00	1237564.0
2018-12-26	1048.34	1066.94	1045.23	1065.00	1237564.0
2018-12-27	1048.34	1066.94	1045.23	1065.00	1237564.0
2018-12-28	1048.34	1066.94	1045.23	1065.00	1237564.0
2018-12-31	1048.34	1066.94	1045.23	1065.00	1237564.0

1255 rows × 5 columns