Database-style Operations on Dataframes

Name: Cuadra, Audrick Zander G.

Section: CPE22S3

Date: March 27, 2024

Submitted to: Engr. Roman Richard

About the data

In this notebook, we will using daily weather data that was taken from the [National Centers for Environmental Information (NCEI) API](). The data collection notebook contains the process that was followed to collect the data.

Note: The NCEI is part of the National Oceanic and Atmospheric Administration (NOAA) and, as you can see from the URL for the API, this resource was created when the NCEI was called the NCDC. Should the URL for this resource change in the future, you can search for the NCEI weather API to find the updated one.

Background on the data

Data meanings:

- `PRCP` : precipitation in millimeters
- `SNOW` : snowfall in millimeters
- `SNWD` : snow depth in millimeters
- `TMAX` : maximum daily temperature in Celsius
- `TMIN` : minimum daily temperature in Celsius
- `TOBS` : temperature at time of observation in Celsius
- `WESF` : water equivalent of snow in millimeters

∨ Setup

```
1 import pandas as p
2
3 weather = p.read_csv('/content/nyc_weather_2018.csv')
4 weather.head()
```

	attributes	datatype	date	station	value
0	,,N,	PRCP	2018-01-01T00:00:00	GHCND:US1CTFR0039	0.0
1	,,N,	PRCP	2018-01-01T00:00:00	GHCND:US1NJBG0015	0.0
2	,,N,	SNOW	2018-01-01T00:00:00	GHCND:US1NJBG0015	0.0
3	,,N,	PRCP	2018-01-01T00:00:00	GHCND:US1NJBG0017	0.0
4	,,N,	SNOW	2018-01-01T00:00:00	GHCND:US1NJBG0017	0.0

- -

Next steps: ⬤ **View recommended plots**

∨ Querying Dataframes

The `query()` method is an easier way of filtering based on some criteria. For example, we can use it to find all entries where snow was recorded:

```
1 snow_data = weather.query('datatype == "SNOW" and value > 0')
2 snow_data.head()
```

	attributes	datatype	date	station	value
124	,,N,	SNOW	2018-01-01T00:00:00	GHCND:US1NYWC0019	25.0
723	,,N,	SNOW	2018-01-04T00:00:00	GHCND:US1NJBG0015	229.0
726	,,N,	SNOW	2018-01-04T00:00:00	GHCND:US1NJBG0017	10.0
730	,,N,	SNOW	2018-01-04T00:00:00	GHCND:US1NJBG0018	46.0
737	,,N,	SNOW	2018-01-04T00:00:00	GHCND:US1NJES0018	10.0

This is equivalent to quering the `data/weather.db` SQLite database for `SELECT * FROM weather WHERE datatype == "SNOW" AND value > 0`:

```
1 import sqlite3
2
3 with sqlite3.connect('/content/weather.db') as connection:
4     snow_data_from_db = p.read_sql(
5         'SELECT * FROM weather WHERE datatype == "SNOW" AND value > 0',
6         connection
7     )
8
9 snow_data.reset_index().drop(columns='index').equals(snow_data_from_db)
```

```
True
```

Note this is also equivalent to creating Boolean masks:

```
1 weather[(weather.datatype == 'SNOW') & (weather.value > 0)].equals(snow_data)
```

```
True
```

Merging DataFrames

We have data for many different stations each day; however, we don't know what the stations are just their IDs. We can join the data in the `data/weather_stations.csv` file which contains information from the `stations` endpoint of the NCEI API. Consult the `weather_data_collection.ipynb` notebook to see how this was collected. It looks like this:

```
1 station_info = p.read_csv('/content/weather_stations.csv')
2 station_info.head()
```

	id	name	latitude	longitude	elevation
0	GHCND:US1CTFR0022	STAMFORD 2.6 SSW, CT US	41.0641	-73.5770	36.6
1	GHCND:US1CTFR0039	STAMFORD 4.2 S, CT US	41.0378	-73.5682	6.4
2	GHCND:US1NJBG0001	BERGENFIELD 0.3 SW, NJ US	40.9213	-74.0020	20.1
	GHCND:US1NJBG0022	SADDLE BROOK TWP			

Next steps: 　🔘 View recommended plots

As a reminder, the weather data looks like this:

```
1 weather.head()
```

	attributes	datatype	date	station	value
0	,,N,	PRCP	2018-01-01T00:00:00	GHCND:US1CTFR0039	0.0
1	,,N,	PRCP	2018-01-01T00:00:00	GHCND:US1NJBG0015	0.0
2	,,N,	SNOW	2018-01-01T00:00:00	GHCND:US1NJBG0015	0.0
3	,,N,	PRCP	2018-01-01T00:00:00	GHCND:US1NJBG0017	0.0
4	,,N,	SNOW	2018-01-01T00:00:00	GHCND:US1NJBG0017	0.0

- -

Next steps: ◖ **View recommended plots**

We can join our data by matching up the `station_info.id` column with the `weather.station` column. Before doing that though, let's see how many unique values we have:

```
1 station_info.id.describe()
```

```
count                       262
unique                      262
top       GHCND:US1CTFR0022
freq                          1
Name: id, dtype: object
```

While `station_info` has one row per station, the `weather` dataframe has many entries per station. Notice it also has fewer uniques:

```
1 weather.station.describe()
```

```
count                    80256
unique                     109
top       GHCND:USW00094789
freq                      4270
Name: station, dtype: object
```

When working with joins, it is important to keep an eye on the row count. Some join types will lead to data loss:

```
1 station_info.shape[0], weather.shape[0]
```

```
(262, 80256)
```

Since we will be doing this often, it makes more sense to write a function:

```
1 def get_row_count(*dfs):
2     return[df.shape[0] for df in dfs]
3 get_row_count(station_info, weather)
```

```
[262, 80256]
```

The **map()** function is more efficient than list comprehensions. We can couple this with **getattr()** to grab any attribute for multiple dataframes:

```
1 def get_info(attr, *dfs):
2     return list(map(lambda x: getattr(x, attr), dfs))
3 get_info('shape', station_info, weather)
```

```
[(262, 5), (80256, 5)]
```

By default `merge()` performs an inner join. We simply specify the columns to use for the join. The left dataframe is the one we call `merge()` on, and the right one is passed in as an argument:

```
1 inner_join = weather.merge(station_info, left_on='station', right_on='id')
2 inner_join.sample(5, random_state=0)
```

	attributes	datatype	date	station	value	
27422	,,N,	PRCP	2018-01-23T00:00:00	GHCND:US1NYSF0061	2.3	GHCND:US1NYS
19317	T,,N,	PRCP	2018-08-10T00:00:00	GHCND:US1NJUN0014	0.0	GHCND:US1NJU
13778	,,N,	WESF	2018-02-18T00:00:00	GHCND:US1NJMS0089	19.6	GHCND:US1NJM
39633	,,7,0700	PRCP	2018-04-06T00:00:00	GHCND:USC00301309	0.0	GHCND:USC003
51025	,,W,2400	SNWD	2018-12-14T00:00:00	GHCND:USW00014734	0.0	GHCND:USW000

We can remove the duplication of information in the `station` and `id` columns by renaming one of them before the merge and then simply using on :

```
1 weather.merge(station_info.rename(dict(id='station'), axis=1), on='station').sample(5, random_state=0)
```

	attributes	datatype	date	station	value	name
27422	,,N,	PRCP	2018-01-23T00:00:00	GHCND:US1NYSF0061	2.3	CENTERPORT 0.9 SW, NY US
19317	T,,N,	PRCP	2018-08-10T00:00:00	GHCND:US1NJUN0014	0.0	WESTFIELD 0.6 NE, NJ US
13778	,,N,	WESF	2018-02-18T00:00:00	GHCND:US1NJMS0089	19.6	PARSIPPANY TROY HILLS TWP 1.3, NJ US
39633	,,7,0700	PRCP	2018-04-06T00:00:00	GHCND:USC00301309	0.0	CENTERPORT, NY US
51025	,,W,2400	SNWD	2018-12-14T00:00:00	GHCND:USW00014734	0.0	NEWARK LIBERTY INTERNATIONAL AIRPORT, NJ US

We are losing stations that don't have weather observations associated with them, if we don't want to lose these rows, we perform a right or left join instead of the inner join:

```
1 left_join = station_info.merge(weather, left_on='id', right_on='station', how='left')
2 right_join = weather.merge(station_info, left_on='station', right_on='id', how='right')
3
4 right_join.tail()
```

	attributes	datatype	date	station	value	
80404	,,W,	WDF5	2018-12-31T00:00:00	GHCND:USW00094789	130.0	GHCND:USW000
80405	,,W,	WSF2	2018-12-31T00:00:00	GHCND:USW00094789	9.8	GHCND:USW000
80406	,,W,	WSF5	2018-12-31T00:00:00	GHCND:USW00094789	12.5	GHCND:USW000
80407	,,W,	WT01	2018-12-31T00:00:00	GHCND:USW00094789	1.0	GHCND:USW000
80408	,,W,	WT02	2018-12-31T00:00:00	GHCND:USW00094789	1.0	GHCND:USW000

The left and right join as we performed above are equivalent because the side that we kept the rows without matches was the same in both cases:

```
1 left_join.sort_index(axis=1).sort_values(['date', 'station']).reset_index().drop(columns='index').equals(
2     right_join.sort_index(axis=1).sort_values(['date', 'station']).reset_index().drop(columns='index')
3 )
```

```
    True
```

Note we have additional rows in the left and right joins because we kept all the stations that didn't have weather observations:

```
1 get_info('shape', inner_join, left_join, right_join)
```

```
    [(80256, 10), (80409, 10), (80409, 10)]
```

If we query the station information for stations that have `NY` in their name, believing that to be all the stations that record weather data for NYC and perform an outer join, we can see where the mismatches occur:

```
1 outer_join = weather.merge(
2     station_info[station_info.name.str.contains('NY')],
3     left_on = 'station', right_on='id', how='outer', indicator=True
4 )
5
6 outer_join.sample(4, random_state=0).append(outer_join[outer_join.station.isna()].head(2))
```

```
<ipython-input-26-291e06d07b42>:6: FutureWarning: The frame.append method is depreca
  outer_join.sample(4, random_state=0).append(outer_join[outer_join.station.isna()].
```

	attributes	datatype	date	station	value	
17259	,,N,	PRCP	2018-05-15T00:00:00	GHCND:US1NJPS0022	0.3	
76178	,,N,	PRCP	2018-05-19T00:00:00	GHCND:US1NJPS0015	8.1	
73410	,,N,	MDPR	2018-08-05T00:00:00	GHCND:US1NYNS0018	12.2	GHCND:US1NYN
74822	,,N,	SNOW	2018-04-02T00:00:00	GHCND:US1NJMS0016	178.0	
80256	NaN	NaN	NaN	NaN	NaN	GHCND:US1NJM
80257	NaN	NaN	NaN	NaN	NaN	GHCND:US1NJM

These joins are equivalent to their SQL counterparts. Below is the inner join. Note that to use `equals()` you will have to do some manipulation of the dataframes to line them up:

```
1 import sqlite3
2
3 with sqlite3.connect('/content/weather.db') as connection:
4     inner_join_from_db = p.read_sql(
5         'SELECT * FROM weather JOIN stations ON weather.station == stations.id',
6         connection
7     )
8
9 inner_join_from_db.shape == inner_join.shape
```

```
    True
```

Revisit the dirty data from the previous module.

```
1 dirty_data = p.read_csv(
2     '/content/dirty_data.csv', index_col='date'
3 ).drop_duplicates().drop(columns='SNWD')
4
5 dirty_data.head()
```

date	station	PRCP	SNOW	TMAX	TMIN	TOBS	WESF	inclement_wea
2018-01-01T00:00:00	?	0.0	0.0	5505.0	-40.0	NaN	NaN	
2018-01-02T00:00:00	GHCND:USC00280907	0.0	0.0	-8.3	-16.1	-12.2	NaN	
2018-01-03T00:00:00	GHCND:USC00280907	0.0	0.0	-4.4	-13.9	-13.3	NaN	
2018-01-								

Next steps: 🔘 **View recommended plots**

We need to create two dataframes for the join. We will drop some unecessary columns as well for easier viewing:

```
1 valid_station = dirty_data.query('station != "?"').copy().drop(columns=['WESF', 'station'])
2 station_with_wesf = dirty_data.query('station == "?"').copy().drop(columns=['station', 'TOBS', 'TMIN', 'TMAX']
```

Our column for the join is the index in both dataframes, so we must specify `left_index` and `right_index` :

```
1 valid_station.merge(
2     station_with_wesf, left_index=True, right_index=True
3 ).query('WESF > 0').head()
```

date	PRCP_x	SNOW_x	TMAX	TMIN	TOBS	inclement_weather_x	PRCP_y	SNOW_y	WI
2018-01-30T00:00:00	0.0	0.0	6.7	-1.7	-0.6	False	1.5	13.0	
2018-03-08T00:00:00	48.8	NaN	1.1	-0.6	1.1	False	28.4	NaN	2
2018-03-13T00:00:00	4.1	51.0	5.6	-3.9	0.0	True	3.0	13.0	
2018-03-21T00:00:00	0.0	0.0	2.8	-2.8	0.6	False	6.6	114.0	
2018-04-02T00:00:00	9.1	127.0	12.8	-1.1	-1.1	True	14.0	152.0	1

Since we are joining on the index, an easier way is to use the `join()` method instead of `merge()` . Note that the suffix parameter is now `lsuffix` for the left dataframe's suffix and `rsuffix` for the right one's:

```
1 valid_station.join(station_with_wesf, rsuffix='_?').query('WESF > 0').head()
```

date	PRCP	SNOW	TMAX	TMIN	TOBS	inclement_weather	PRCP_?	SNOW_?	WESF	i
2018-01-30T00:00:00	0.0	0.0	6.7	-1.7	-0.6	False	1.5	13.0	1.8	
2018-03-08T00:00:00	48.8	NaN	1.1	-0.6	1.1	False	28.4	NaN	28.7	
2018-03-13T00:00:00	4.1	51.0	5.6	-3.9	0.0	True	3.0	13.0	3.0	
2018-03-						F		114.0		

Joins can be very resource-intensive, so it's a good idea to figure out what type of join you need using set operations before trying the join itself. The pandas set operations are performed on the index, so whichever columns we will be joining on will need to be the index. Let's go back to the `weather` and `station_info` dataframes and set the station ID columns as the index:

```
1 weather.set_index('station', inplace=True)
2 station_info.set_index('id', inplace=True)
```

The intersection will tell us the stations that are present in both dataframes. The result will be the index when performing an inner join:

```
1 weather.index.intersection(station_info.index)
```

```
Index(['GHCND:US1CTFR0039', 'GHCND:US1NJBG0015', 'GHCND:US1NJBG0017',
       'GHCND:US1NJBG0018', 'GHCND:US1NJBG0023', 'GHCND:US1NJBG0030',
       'GHCND:US1NJBG0039', 'GHCND:US1NJBG0044', 'GHCND:US1NJES0018',
```

```
            'GHCND:US1NJES0024',
            ...
            'GHCND:US1NJMS0047', 'GHCND:US1NYSF0083', 'GHCND:US1NYNY0074',
            'GHCND:US1NJPS0018', 'GHCND:US1NJBG0037', 'GHCND:USC00284987',
            'GHCND:US1NJES0031', 'GHCND:US1NJMD0086', 'GHCND:US1NJMS0097',
            'GHCND:US1NJMN0081'],
          dtype='object', length=109)
```

The set difference will tell us what we lose from each side. When performing an inner join, we lose nothing from the `weather` dataframe:

```
1 weather.index.difference(station_info.index)
```

```
    Index([], dtype='object')
```

We lose 153 stations from the `station_info` dataframe, however:

```
1 station_info.index.difference(weather.index)
```

```
    Index(['GHCND:US1CTFR0022', 'GHCND:US1NJBG0001', 'GHCND:US1NJBG0002',
           'GHCND:US1NJBG0005', 'GHCND:US1NJBG0006', 'GHCND:US1NJBG0008',
           'GHCND:US1NJBG0011', 'GHCND:US1NJBG0012', 'GHCND:US1NJBG0013',
           'GHCND:US1NJBG0020',
           ...
           'GHCND:USC00308322', 'GHCND:USC00308749', 'GHCND:USC00308946',
           'GHCND:USC00309117', 'GHCND:USC00309270', 'GHCND:USC00309400',
           'GHCND:USC00309466', 'GHCND:USC00309576', 'GHCND:USW00014708',
           'GHCND:USW00014786'],
          dtype='object', length=153)
```

The symmetric difference will tell us what gets lost from both sides. It is the combination of the set difference in both directions:

```
1 ny_in_name = station_info[station_info.name.str.contains('NY')]
2
3 ny_in_name.index.difference(weather.index).shape[0]\
4 + weather.index.difference(ny_in_name.index).shape[0]\
5 == weather.index.symmetric_difference(ny_in_name.index).shape[0]
```

```
    True
```

The union will show us everything that will be present after a full outer join. Note that since these are sets (which don't allow duplicates by definition), we must pass unique entries for union:

```
1 weather.index.unique().union(station_info.index)
```

```
    Index(['GHCND:US1CTFR0022', 'GHCND:US1CTFR0039', 'GHCND:US1NJBG0001',
           'GHCND:US1NJBG0002', 'GHCND:US1NJBG0003', 'GHCND:US1NJBG0005',
           'GHCND:US1NJBG0006', 'GHCND:US1NJBG0008', 'GHCND:US1NJBG0010',
           'GHCND:US1NJBG0011',
           ...
           'GHCND:USW00014708', 'GHCND:USW00014732', 'GHCND:USW00014734',
           'GHCND:USW00014786', 'GHCND:USW00054743', 'GHCND:USW00054787',
           'GHCND:USW00094728', 'GHCND:USW00094741', 'GHCND:USW00094745',
           'GHCND:USW00094789'],
          dtype='object', length=262)
```

Note that the symmetric difference is actually the union of the set differences:

```
1 ny_in_name = station_info[station_info.name.str.contains('NY')]
2
3 ny_in_name.index.difference(weather.index).union(weather.index.difference(ny_in_name.index)).equals(
4     weather.index.symmetric_difference(ny_in_name.index)
```

```
5 )
```

True