Aggregations with pandas and numpy

Name: Cuadra, Audrick Zander G.

Section: CPE22S3

Date: March 27, 2024

Submitted to: Engr. Roman Richard

About the data

In this notebook, we will be working with 2 data sets:

- Facebook's stock price throughout 2018 (obtained using the `stock_analysis` [package](#)).
- daily weather data for NYC from the [National Centers for Environmental Information (NCEI) API](#).

Note: The NCEI is part of the National Oceanic and Atmospheric Administration (NOAA) and, as you can see from the URL for the API, this resource was created when the NCEI was called the NCDC. Should the URL for this resource change in the future, you can search for the NCEI weather API to find the updated one.

Background on the weather data

Data meanings:

- `AWND` : average wind speed
- `PRCP` : precipitation in millimeters
- `SNOW` : snowfall in millimeters
- `SNWD` : snow depth in millimeters
- `TMAX` : maximum daily temperature in Celsius
- `TMIN` : minimum daily temperature in Celsius

⌄ Setup

```
1 import numpy as ny
2 import pandas as p
3
4 weather = p.read_csv('/content/weather_by_station.csv', index_col='date', parse_dates=
5 weather.head()
```

	datatype	station	value	station_name
date				
2018-01-01	PRCP	GHCND:US1CTFR0039	0.0	STAMFORD 4.2 S, CT US
2018-01-01	PRCP	GHCND:US1NJBG0015	0.0	NORTH ARLINGTON 0.7 WNW, NJ US
2018-01-01	SNOW	GHCND:US1NJBG0015	0.0	NORTH ARLINGTON 0.7 WNW, NJ US
2018-01-01	PRCP	GHCND:US1NJBG0017	0.0	GLEN ROCK 0.7 SSE, NJ US
2018-01-01	SNOW	GHCND:US1NJBG0017	0.0	GLEN ROCK 0.7 SSE, NJ US

```
1 fb = p.read_csv('/content/fb_2018.csv', index_col='date', parse_dates=True).assign(
2   trading_volume=lambda x: p.cut(x.volume, bins=3, labels=['low', 'med', 'high'])
3 )
4 fb.head()
```

	open	high	low	close	volume	trading_volume
date						
2018-01-02	177.68	181.58	177.5500	181.42	18151903	low
2018-01-03	181.88	184.78	181.3300	184.67	16886563	low
2018-01-04	184.90	186.21	184.0996	184.33	13880896	low
2018-01-05	185.59	186.90	184.9300	186.85	13574535	low
2018-01-08	187.20	188.90	186.3300	188.28	17994726	low

Before we dive into any calculations, let's make sure pandas won't put things in scientific notation. We will modify how floats are formatted for displaying. The format we will apply is `.2f`, which will provide the float with 2 digits after the decimal point:

```
1 p.set_option('display.float_format', lambda x: '%.2f' % x)
```

⌄ Summarizing DataFrames

We learned about agg() in the dataframe operations notebook when we learned about window calculations; however, we can call this on the dataframe directly to aggregate its contents into a single series:

```
1 fb.agg({
2     'open': ny.mean,
3     'high': ny.max,
4     'low': ny.min,
5     'close': ny.mean,
6     'volume': ny.sum
7 })
```

```
open                 171.45
high                 218.62
low                  123.02
close                171.51
volume     6949682394.00
dtype: float64
```

We can use this to find the total snowfall and precipitation recorded in Central Park in 2018:

```
1 weather.query(
2     'station == "GHCND:USW00094728"'
3 ).pivot(columns='datatype', values='value')[['SNOW', 'PRCP']].sum()
```

```
datatype
SNOW    1007.00
PRCP    1665.30
dtype: float64
```

This is equivalent to passing `'sum'` to `agg()`:

```
1 weather.query(
2     'station == "GHCND:USW00094728"'
3 ).pivot(columns='datatype', values='value')[['SNOW', 'PRCP']].agg('sum')
```

```
datatype
SNOW    1007.00
PRCP    1665.30
dtype: float64
```

Note that we aren't limited to providing a single aggregation per column. We can pass a list, and we will get a dataframe back instead of a series. nan values are placed where we don't have a calculation result to display:

```
1 fb.agg({
2     'open': 'mean',
3     'high': ['min', 'max'],
4     'low': ['min', 'max'],
5     'close': 'mean'
6 })
```

	open	high	low	close
mean	171.45	NaN	NaN	171.51
min	NaN	129.74	123.02	NaN
max	NaN	218.62	214.27	NaN

⌄ Using `groupby()`

Often we won't want to aggregate on the entire dataframe, but on groups within it. For this purpose, we can run `groupby()` before the aggregation. If we group by the `trading_volume` column, we will get a row for each of the values it takes on:

```
1 fb.groupby('trading_volume').mean()
```

	open	high	low	close	volume
trading_volume					
low	171.36	173.46	169.31	171.43	24547207.71
med	175.82	179.42	172.11	175.14	79072559.12
high	167.73	170.48	161.57	168.16	141924023.33

After we run the `groupby()`, we can still select columns for aggregation

```
1 fb.groupby('trading_volume')['close'].agg(['min', 'max', 'mean'])
```

	min	max	mean
trading_volume			
low	124.06	214.67	171.43
med	152.22	217.50	175.14
high	160.06	176.26	168.16

We can still provide a dictionary specifying the aggregations to perform, but passing a list for a column will result in a hierarchical index for the columns:

```
1 fb_agg = fb.groupby('trading_volume').agg({
2     'open': 'mean',
3     'high': ['min', 'max'],
4     'low': ['min', 'max'],
5     'close': 'mean'
6 })
7 fb_agg
```

	open	high		low		close
	mean	min	max	min	max	mean
trading_volume						
low	171.36	129.74	216.20	123.02	212.60	171.43
med	175.82	162.85	218.62	150.75	214.27	175.14
high	167.73	161.10	180.13	149.02	173.75	168.16

The hierarchical index in the columns looks like this:

```
1 fb_agg.columns
```

```
MultiIndex([( 'open', 'mean'),
            ( 'high',  'min'),
            ( 'high',  'max'),
            (  'low',  'min'),
            (  'low',  'max'),
            ('close', 'mean')],
           )
```

Using a list comprehension, we can join the levels (in a tuple) with an _ at each iteration:

```
1 fb_agg.columns = ['_'.join(col_agg) for col_agg in fb_agg.columns]
2 fb_agg.head()
```

	open_mean	high_min	high_max	low_min	low_max	close_mean
trading_volume						
low	171.36	129.74	216.20	123.02	212.60	171.43
med	175.82	162.85	218.62	150.75	214.27	175.14
high	167.73	161.10	180.13	149.02	173.75	168.16

We can group on datetimes despite them being in the index if we use a `Grouper`:

```
1 weather['2018-10'].query('datatype == "PRCP"').groupby(
2     p.Grouper(freq='D')
3 ).mean().head()
```

```
<ipython-input-17-927d8e720cd7>:1: FutureWarning: Indexing a DataFrame with a datetim
  weather['2018-10'].query('datatype == "PRCP"').groupby(
<ipython-input-17-927d8e720cd7>:3: FutureWarning: The default value of numeric_only i
  ).mean().head()
```

	value
date	
2018-10-01	0.01
2018-10-02	2.23
2018-10-03	19.69
2018-10-04	0.32
2018-10-05	0.97

This `Grouper` can be one of many group by values. Here, we find the quarterly total precipitation per station:

```
1 weather.query('datatype == "PRCP"').groupby(
2     ['station_name', p.Grouper(freq='Q')]
3 ).sum().unstack().sample(5, random_state=1)
```

```
<ipython-input-20-648f26f48c8d>:3: FutureWarning: The default value of numeric_only i
  ).sum().unstack().sample(5, random_state=1)
```

	value			
date	2018-03-31	2018-06-30	2018-09-30	2018-12-31
station_name				
WANTAGH 1.1 NNE, NY US	279.90	216.80	472.50	277.20
STATEN ISLAND 1.4 SE, NY US	379.40	295.30	438.80	409.90
SYOSSET 2.0 SSW, NY US	323.50	263.30	355.50	459.90
STAMFORD 4.2 S, CT US	338.00	272.10	424.70	390.00
WAYNE TWP 0.8 SSW, NJ US	246.20	295.30	620.90	422.00

Note that we can use `filter()` to exclude some groups from aggregation. Here, we only keep groups with 'NY' in the group's `name` attribute, which is the station ID in this case:

```
1 weather.groupby('station').filter( # station IDs with NY in them
2     lambda x: 'NY' in x.name
3 ).query('datatype == "SNOW"').groupby('station_name').sum().squeeze() # aggregate and
```

```
<ipython-input-22-799de504673b>:3: FutureWarning: The default value of numeric_only i
  ).query('datatype == "SNOW"').groupby('station_name').sum().squeeze() # aggregate a
```

```
station_name
ALBERTSON 0.2 SSE, NY US              1087.00
AMITYVILLE 0.1 WSW, NY US              434.00
AMITYVILLE 0.6 NNE, NY US             1072.00
ARMONK 0.3 SE, NY US                  1504.00
BROOKLYN 3.1 NW, NY US                 305.00
CENTERPORT 0.9 SW, NY US               799.00
ELMSFORD 0.8 SSW, NY US                863.00
FLORAL PARK 0.4 W, NY US              1015.00
HICKSVILLE 1.3 ENE, NY US              716.00
JACKSON HEIGHTS 0.3 WSW, NY US         107.00
LOCUST VALLEY 0.3 E, NY US               0.00
LYNBROOK 0.3 NW, NY US                 325.00
MASSAPEQUA 0.9 SSW, NY US               41.00
MIDDLE VILLAGE 0.5 SW, NY US          1249.00
NEW HYDE PARK 1.6 NE, NY US              0.00
NEW YORK 8.8 N, NY US                    0.00
NORTH WANTAGH 0.4 WSW, NY US           471.00
PLAINEDGE 0.4 WSW, NY US               610.00
PLAINVIEW 0.4 ENE, NY US              1360.00
SADDLE ROCK 3.4 WSW, NY US             707.00
STATEN ISLAND 1.4 SE, NY US            936.00
STATEN ISLAND 4.5 SSE, NY US            89.00
SYOSSET 2.0 SSW, NY US                1039.00
VALLEY STREAM 0.6 SE, NY US            898.00
WANTAGH 0.3 ESE, NY US                1280.00
WANTAGH 1.1 NNE, NY US                 940.00
WEST NYACK 1.3 WSW, NY US             1371.00
Name: value, dtype: float64
```

Let's see which months have the most precipitation. First, we need to group by day and average the precipitation across the stations. Then we can group by month and sum the resulting precipitation. We use `nlargest()` to give the 5 months with the most precipitation:

```
1 weather.query('datatype == "PRCP"').groupby(
2     p.Grouper(freq='D')
3 ).mean().groupby(p.Grouper(freq='M')).sum().value.nlargest()
```

```
<ipython-input-23-b52d50d03a1a>:3: FutureWarning: The default value of numeric_only i
  ).mean().groupby(p.Grouper(freq='M')).sum().value.nlargest()
date
2018-11-30    210.59
2018-09-30    193.09
2018-08-31    192.45
2018-07-31    160.98
2018-02-28    158.11
Name: value, dtype: float64
```

Perhaps the previous result was surprising. The saying goes "April showers bring May flowers"; yet April wasn't in the top 5 (neither was May for that matter). Snow will count towards precipitation, but that doesn't explain why summer months are higher than April. Let's look for days that accounted for a large percentage of the precipitation in a given month.

In order to do so, we need to calculate the average daily precipitation across stations and then find the total per month. This will be the denominator. However, in order to divide the daily values by the total for their month, we will need a Series of equal dimensions. This means we will need to use `transform()`:

```
1 weather.query('datatype == "PRCP"').rename(
2     dict(value='prcp'), axis=1
3 ).groupby(p.Grouper(freq='D')).mean().groupby(
4     p.Grouper(freq='M')
5 ).transform(ny.sum)['2018-01-28':'2018-02-03']
```

```
<ipython-input-24-1e336928755e>:3: FutureWarning: The default value of numeric_only i
  ).groupby(p.Grouper(freq='D')).mean().groupby(
```

	prcp
date	
2018-01-28	69.31
2018-01-29	69.31
2018-01-30	69.31
2018-01-31	69.31
2018-02-01	158.11
2018-02-02	158.11
2018-02-03	158.11

Notice how we have the same value repeated for each day in the month it belongs to. This will allow us to calculate the percentage of the monthly precipitation that occurred each day and then pull out the largest values:

```
1 weather\
2     .query('datatype == "PRCP"')\
3     .rename(dict(value='prcp'), axis=1)\
4     .groupby(p.Grouper(freq='D')).mean()\
5     .assign(
6     total_prcp_in_month=lambda x: x.groupby(
7         p.Grouper(freq='M')
8     ).transform(ny.sum),
9     pct_monthly_prcp=lambda x: x.prcp.div(
10        x.total_prcp_in_month
11     )
12     ).nlargest(5, 'pct_monthly_prcp')
```

```
<ipython-input-26-ec955deab43a>:4: FutureWarning: The default value of numeric_only i
  .groupby(p.Grouper(freq='D')).mean()\
```

date	prcp	total_prcp_in_month	pct_monthly_prcp
2018-10-12	34.77	105.63	0.33
2018-01-13	21.66	69.31	0.31
2018-03-02	38.77	137.46	0.28
2018-04-16	39.34	140.57	0.28
2018-04-17	37.30	140.57	0.27

transform() can be used on dataframes as well. We can use it to easily standardize the data:

```
1 fb[['open', 'high', 'low', 'close']].transform(
2     lambda x: (x - x.mean()).div(x.std())
3 ).head()
```

date	open	high	low	close
2018-01-02	0.32	0.41	0.41	0.50
2018-01-03	0.53	0.57	0.60	0.66
2018-01-04	0.68	0.65	0.74	0.64
2018-01-05	0.72	0.68	0.78	0.77
2018-01-08	0.80	0.79	0.85	0.84

∨ Pivot tables and crosstabs

We saw pivots in before; however, we weren't able to provide any aggregations. With `pivot_table()`, we get the mean by default as the `aggfunc`. In its simplest form, we provide a column to place along the columns:

```
1 fb.pivot_table(columns='trading_volume')
```

trading_volume	low	med	high
close	171.43	175.14	168.16
high	173.46	179.42	170.48
low	169.31	172.11	161.57
open	171.36	175.82	167.73
volume	24547207.71	79072559.12	141924023.33

By placing the trading volume in the index, we get the aggregation from the first example in the group by section above:

```
1 fb.pivot_table(index='trading_volume')
```

	close	high	low	open	volume
trading_volume					
low	171.43	173.46	169.31	171.36	24547207.71
med	175.14	179.42	172.11	175.82	79072559.12
high	168.16	170.48	161.57	167.73	141924023.33

With `pivot()`, we also weren't able to handle multi-level indices or indices with repeated values. For this reason we haven't been able to put the weather data in the wide format. The `pivot_table()` method solves this issue:

```
1 weather.reset_index().pivot_table(
2     index=['date', 'station', 'station_name'],
3     columns='datatype',
4     values='value',
5     aggfunc='median'
6 ).reset_index().tail()
```

	datatype	date	station	station_name	AWND	DAPR	MDPR	PGTM	PRCP
28740		2018-12-31	GHCND:USW00054787	FARMINGDALE REPUBLIC AIRPORT, NY US	5.00	NaN	NaN	2052.00	28.70
28741		2018-12-31	GHCND:USW00094728	NY CITY CENTRAL PARK, NY US	NaN	NaN	NaN	NaN	25.90
28742		2018-12-31	GHCND:USW00094741	TETERBORO AIRPORT, NJ US	1.70	NaN	NaN	1954.00	29.20
28743		2018-12-31	GHCND:USW00094745	WESTCHESTER CO AIRPORT, NY US	2.70	NaN	NaN	2212.00	24.40
28744		2018-12-31	GHCND:USW00094789	JFK INTERNATIONAL AIRPORT, NY US	4.10	NaN	NaN	NaN	31.20

5 rows × 30 columns

We can use the `pd.crosstab()` function to create a frequency table. For example, if we want to see how many low-, medium-, and high-volume trading days Facebook

```
1 p.crosstab(
2     index=fb.trading_volume,
3     columns=fb.index.month,
4     colnames=['month'] # name the columns index
5 )
```

month	1	2	3	4	5	6	7	8	9	10	11	12
trading_volume												
low	20	19	15	20	22	21	18	23	19	23	21	19
med	1	0	4	1	0	0	2	0	0	0	0	0
high	0	0	2	0	0	0	1	0	0	0	0	0

We can normalize with the row or column totals with the `normalize` parameter. This shows percentage of the total:

```
1 p.crosstab(
2     index=fb.trading_volume,
3     columns=fb.index.month,
4     colnames=['month'],
5     normalize='columns'
6 )
```

trading_volume	month	1	2	3	4	5	6	7	8	9	10	11	12
low		0.95	1.00	0.71	0.95	1.00	1.00	0.86	1.00	1.00	1.00	1.00	1.00
med		0.05	0.00	0.19	0.05	0.00	0.00	0.10	0.00	0.00	0.00	0.00	0.00
high		0.00	0.00	0.10	0.00	0.00	0.00	0.05	0.00	0.00	0.00	0.00	0.00

If we want to perform a calculation other than counting the frequency, we can pass the column to run the calculation on to `values` and the function to use to `aggfunc`:

```
1 p.crosstab(
2     index=fb.trading_volume,
3     columns=fb.index.month,
4     colnames=['month'],
5     values=fb.close,
6     aggfunc=ny.mean
7 )
```

trading_volume	month	1	2	3	4	5	6	7	8	9
low		185.24	180.27	177.07	163.29	182.93	195.27	201.92	177.49	164.38	1!
med		179.37	NaN	164.76	174.16	NaN	NaN	194.28	NaN	NaN
high		NaN	NaN	164.11	NaN	NaN	NaN	176.26	NaN	NaN

We can also get row and column subtotals with the `margins` parameter. Let's count the number of times each station recorded snow per month and include the subtotals:

```
 1 snow_data = weather.query('datatype == "SNOW"')
 2 p.crosstab(
 3     index=snow_data.station_name,
 4     columns=snow_data.index.month,
 5     colnames=['month'],
 6     values=snow_data.value,
 7     aggfunc=lambda x: (x > 0).sum(),
 8     margins=True, # show row and column subtotals
 9     margins_name='total observations of snow' # name the subtotals
10 )
```

month	1	2	3	4	5	6	7	8	9	10	11
station_name											
ALBERTSON 0.2 SSE, NY US	3.00	1.00	3.00	1.00	0.00	0.00	0.00	0.00	0.00	0.00	1.00
AMITYVILLE 0.1 WSW, NY US	1.00	0.00	1.00	1.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
AMITYVILLE 0.6 NNE, NY US	3.00	1.00	3.00	1.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
ARMONK 0.3 SE, NY US	6.00	4.00	6.00	3.00	0.00	0.00	0.00	0.00	0.00	0.00	1.00
BLOOMINGDALE 0.7 SSE, NJ US	2.00	1.00	3.00	1.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00
...	...	...	...	...	...	...	...	...	...	...	...
WESTFIELD 0.6 NE, NJ US	3.00	0.00	4.00	1.00	0.00	NaN	0.00	0.00	0.00	NaN	1.00
WOODBRIDGE											

⌄ Comments and Insights

The module topic demonstrates Aggregations. The module displays proper techniques in summarizing and grouping data with the use of different functions and module attribute which makes data presentation look more neat and organize.